

Predicting the Quarterly Revenue for Alibaba

```
In [1]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [2]: 1 # Reading the Data
2 df=pd.read_excel('Alibaba Quarterly Revenue.xlsx')
3 df.head()
```

```
Out[2]:
```

	Date	Quarterly Revenue
0	2013-09-30	1777
1	2013-12-31	3061
2	2014-06-30	2542
3	2014-09-30	2742
4	2014-12-31	4219

```
In [3]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36 entries, 0 to 35
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0    Date            36 non-null    datetime64[ns]
1    Quarterly Revenue  36 non-null    int64  
dtypes: datetime64[ns](1), int64(1)
memory usage: 704.0 bytes
```

```
In [4]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

```
Out[4]:
```

	Quarterly Revenue
Date	
2013-09-30	1777
2013-12-31	3061
2014-06-30	2542
2014-09-30	2742
2014-12-31	4219

```
In [5]: 1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Alibaba Quarterly Revenue')
4 plt.show()
```




```
In [8]: 1 # Dividing the data into training and testing
2 # Plotting the data
3 import seaborn as sns
4 df['Date'] = df.index
5 train = df[df['Date'] < pd.to_datetime("2020-12", format='%Y-%m')]
6 train['train'] = train['Quarterly Revenue']
7 del train['Date']
8 del train['Quarterly Revenue']
9 test = df[df['Date'] >= pd.to_datetime("2020-12", format='%Y-%m')]
10 del test['Date']
11 test['test'] = test['Quarterly Revenue']
12 del test['Quarterly Revenue']
13 plt.plot(train, color = "black")
14 plt.plot(test, color = "red")
15 plt.title("Train/Test split for Quarterly Revenue")
16 plt.ylabel("Quarterly Revenue")
17 plt.xlabel('Date')
18 sns.set()
19 plt.show()
```

C:\Users\ravit\AppData\Local\Temp\ipykernel_28024\1521455262.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
train['train'] = train['Quarterly Revenue']

C:\Users\ravit\AppData\Local\Temp\ipykernel_28024\1521455262.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
test['test'] = test['Quarterly Revenue']



Arima Model

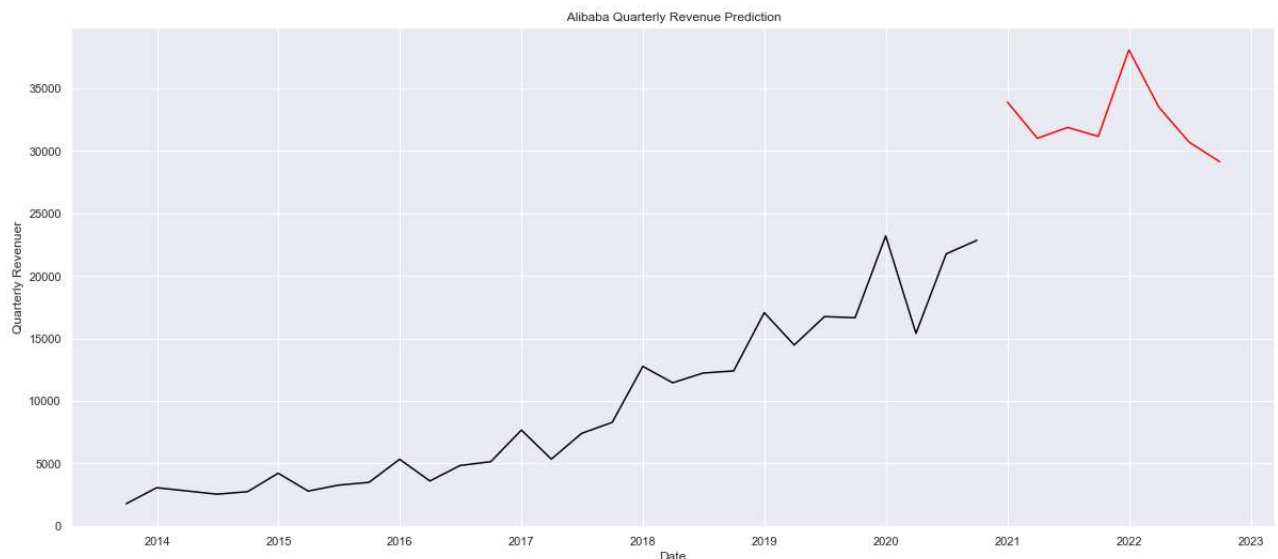
```
In [9]: 1 # Applying ARIMA Model
2 from pmdarima.arima import auto_arima
3 model = auto_arima(train, trace=True, error_action='ignore', suppress_warnings=True)
4 model.fit(train)
5 forecast = model.predict(n_periods=len(test))
6 forecast = pd.DataFrame(forecast, index = test.index, columns=['Prediction'])
```

```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=496.559, Time=0.39 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=509.779, Time=0.01 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=495.750, Time=0.02 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=496.533, Time=0.12 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=509.765, Time=0.01 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=495.450, Time=0.03 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=491.068, Time=0.06 sec
ARIMA(4,1,0)(0,0,0)[0] intercept : AIC=inf, Time=0.22 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=488.303, Time=0.31 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=496.118, Time=0.24 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=485.002, Time=0.40 sec
ARIMA(5,1,1)(0,0,0)[0] intercept : AIC=486.920, Time=0.42 sec
ARIMA(4,1,2)(0,0,0)[0] intercept : AIC=487.007, Time=0.34 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=487.113, Time=0.33 sec
ARIMA(5,1,0)(0,0,0)[0] intercept : AIC=inf, Time=0.36 sec
ARIMA(5,1,2)(0,0,0)[0] intercept : AIC=488.008, Time=0.51 sec
ARIMA(4,1,1)(0,0,0)[0] : AIC=484.662, Time=0.26 sec
ARIMA(3,1,1)(0,0,0)[0] : AIC=500.488, Time=0.21 sec
ARIMA(4,1,0)(0,0,0)[0] : AIC=inf, Time=0.20 sec
ARIMA(5,1,1)(0,0,0)[0] : AIC=486.554, Time=0.35 sec
ARIMA(4,1,2)(0,0,0)[0] : AIC=486.708, Time=0.16 sec
ARIMA(3,1,0)(0,0,0)[0] : AIC=506.117, Time=0.04 sec
ARIMA(3,1,2)(0,0,0)[0] : AIC=inf, Time=0.22 sec
ARIMA(5,1,0)(0,0,0)[0] : AIC=486.141, Time=0.12 sec
ARIMA(5,1,2)(0,0,0)[0] : AIC=488.707, Time=0.30 sec
```

```
Best model: ARIMA(4,1,1)(0,0,0)[0]
Total fit time: 5.663 seconds
```

```
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:834: ValueWarning: No supported index is available.
Prediction results will be given with an integer index beginning at `start`.
    return get_prediction_index()
```

```
In [10]: 1 # Plotting the prediction
2 plt.plot(train, color = "black")
3 plt.plot(test, color = "red")
4 plt.plot(forecast, color = "green")
5 plt.title(" Alibaba Quarterly Revenue Prediction")
6 plt.ylabel("Quarterly Revenuer")
7 plt.xlabel('Date')
8 sns.set()
9 plt.show()
```



SARIMA Model

```
In [22]: 1 df=pd.read_excel('Alibaba Quarterly Revenue.xlsx')
2 df = df.set_index('Date')
```

```
In [23]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)

SARIMAX: (0, 0, 1) x (0, 1, 0, 12)

SARIMAX: (0, 1, 0) x (0, 1, 1, 12)

SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [24]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_trend = False)
6             result = mod.fit()
7             print('SARIMA{}x{}12 - AIC:{}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

```
In [26]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                   seasonal_order = (1, 1, 0, 12))
5
6 result = model.fit()
7 print(result.summary().tables[1])
```

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

If an index is available, see if it is a date-based index or if it

```
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1         -0.5736      0.586     -0.978      0.328     -1.723      0.576
ma.L1          0.4940      0.632      0.781      0.435     -0.745      1.733
ar.S.L12        0.0019      0.164      0.011      0.991     -0.320      0.324
sigma2         6.155e+06    2.96e+08    2.08e+14    0.000    6.15e+06    6.15e+06
=====
```


```
In [72]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-04-30'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/base/tsa_model.py:132: FutureWarning: The 'freq' argument in Timestamp is deprecated and will be removed in a future version.
Negative indices (that lie in the Index)

```
Out[72]:
```

	lower Quarterly Revenue	upper Quarterly Revenue
2021-04-30	133299.243434	143023.993307
2021-07-31	138972.339021	148697.088895
2021-10-31	133165.550096	142890.299970
2022-01-31	149300.397194	159025.147067
2022-04-30	133392.899885	143117.649759
2022-07-31	142751.969848	152476.719721
2022-10-31	145428.832134	155153.582008


```
In [73]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Revenue')
7 plt.legend()
8 plt.show()
```



```
In [74]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-04-30:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 8648353.268553345
The Root Mean Squared Error of our forecasts is 2940.808267900739

```
In [38]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Revenue')
10
11 plt.legend()
12 plt.show()
13
```



DNN MODEL

```
In [39]: 1 def convert2matrix(data_arr, look_back):
2     X, Y = [], []
3     for i in range(len(data_arr)-look_back):
4         d=i+look_back
5         X.append(data_arr[i:d,0])
6         Y.append(data_arr[d,0])
7     return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [43]: 1 df=pd.read_excel('Alibaba Quarterly Revenue.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [44]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup Look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [45]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [46]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)

Epoch 1/500
12/12 [=====] - 1s 15ms/step - loss: 6500924416.0000 - mse: 6500924416.0000 - mae: 80430.8203 - val_loss: 7176567296.0000 - val_mse: 7176567296.0000 - val_mae: 84625.9297
Epoch 2/500
12/12 [=====] - 0s 3ms/step - loss: 3451995136.0000 - mse: 3451995136.0000 - mae: 58461.0117 - val_loss: 3339458816.0000 - val_mse: 3339458816.0000 - val_mae: 57673.1992
Epoch 3/500
12/12 [=====] - 0s 3ms/step - loss: 1369991680.0000 - mse: 1369991680.0000 - mae: 36457.2852 - val_loss: 1036054336.0000 - val_mse: 1036054336.0000 - val_mae: 32019.2207
Epoch 4/500
12/12 [=====] - 0s 3ms/step - loss: 326478656.0000 - mse: 326478656.0000 - mae: 16779.9941 - val_loss: 140287584.0000 - val_mse: 140287584.0000 - val_mae: 11434.0312
Epoch 5/500
12/12 [=====] - 0s 3ms/step - loss: 49302148.0000 - mse: 49302148.0000 - mae: 5308.9678 - val_loss: 9422529.0000 - val_mse: 9422529.0000 - val_mae: 2774.2708
Epoch 6/500
12/12 [=====] - 0s 4ms/step - loss: 49988408.0000 - mse: 49988408.0000 - mae: 6629.2236 - val_loss: 19770046.0000 - val_mse: 19770046.0000 - val_mae: 3221.9114
Epoch 7/500
12/12 [=====] - 0s 4ms/step - loss: 49988408.0000 - mse: 49988408.0000 - mae: 6629.2236 - val_loss: 19770046.0000 - val_mse: 19770046.0000 - val_mae: 3221.9114
```

```
In [47]: 1 def model_loss(history):
2     plt.figure(figsize=(8,4))
3     plt.plot(history.history['loss'], label='Train Loss')
4     plt.plot(history.history['val_loss'], label='Test Loss')
5     plt.title('model loss')
6     plt.ylabel('loss')
7     plt.xlabel('epochs')
8     plt.legend(loc='upper right')
9     plt.show();
```

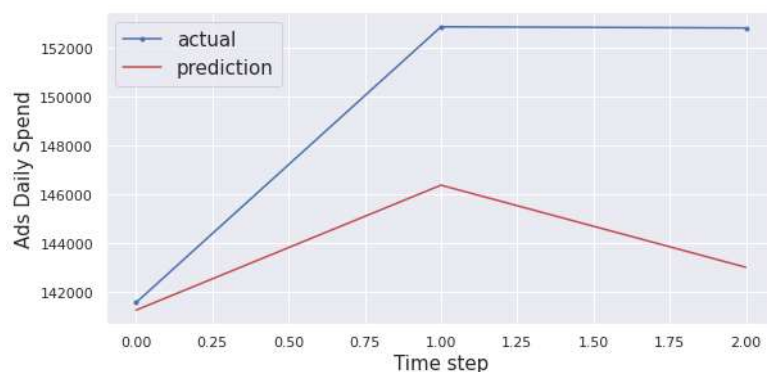
```
In [48]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))

Train Root Mean Squared Error(RMSE): 2753.37; Train Mean Absolute Error(MAE) : 2053.85
[7581066.5, 7581066.5, 2053.8525390625]
[46036500.0, 46036500.0, 5528.796875]
Test Root Mean Squared Error(RMSE): 6785.02; Test Mean Absolute Error(MAE) : 5528.80
```

```
In [49]: 1 def prediction_plot(testY, test_predict):
2     len_prediction=[x for x in range(len(testY))]
3     plt.figure(figsize=(8,4))
4     plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5     plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6     plt.tight_layout()
7     sns.despine(top=True)
8     plt.subplots_adjust(left=0.07)
9     plt.ylabel('Ads Daily Spend', size=15)
10    plt.xlabel('Time step', size=15)
11    plt.legend(fontsize=15)
12    plt.show();
```

```
In [50]: 1 test_predict = model.predict(testX)
2 prediction_plot(testY, test_predict)
```

1/1 [=====] - 0s 76ms/step



GRU and BiLSTM Models

```
In [51]: 1 df=pd.read_excel('Alibaba Quarterly Revenue.xlsx')
        2 df.head()
```

```
Out[51]:
```

	Date	Quarterly Revenue
0	2009-01-31	108627
1	2009-04-30	94242
2	2009-07-31	100876
3	2009-10-31	99373
4	2010-01-31	113594

```
In [52]: 1 df = df.set_index('Date')
        2 df.head()
```

```
Out[52]:
```

	Date	Quarterly Revenue
	2009-01-31	108627
	2009-04-30	94242
	2009-07-31	100876
	2009-10-31	99373
	2010-01-31	113594

```
In [53]: 1 # Split train data and test data
        2 train_size = int(len(df)*0.8)
        3
        4 train_data = df.iloc[:train_size]
        5 test_data = df.iloc[train_size:]
```

```
In [54]: 1 scaler = MinMaxScaler().fit(train_data)
        2 train_scaled = scaler.transform(train_data)
        3 test_scaled = scaler.transform(test_data)
```

```
In [55]: 1 # Create input dataset
        2 def create_dataset (X, look_back = 1):
        3     Xs, ys = [], []
        4
        5     for i in range(len(X)-look_back):
        6         v = X[i:i+look_back]
        7         Xs.append(v)
        8         ys.append(X[i+look_back])
        9
        10    return np.array(Xs), np.array(ys)
        11 LOOK_BACK = 4
        12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
        13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
        14 # Print data shape
        15 print('X_train.shape: ', X_train.shape)
        16 print('y_train.shape: ', y_train.shape)
        17 print('X_test.shape: ', X_test.shape)
        18 print('y_test.shape: ', y_test.shape)
```

```
X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (8, 4, 1)
y_test.shape: (8, 1)
```



```
In [56]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```



```
In [65]: 1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3         patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5         validation_split = 0.2,
6         batch_size = 16, shuffle = False,
7         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```

```
Epoch 1/100
2/2 [=====] - 0s 85ms/step - loss: 0.0187 - val_loss: 0.0225
Epoch 2/100
2/2 [=====] - 0s 80ms/step - loss: 0.0165 - val_loss: 0.0200
Epoch 3/100
2/2 [=====] - 0s 44ms/step - loss: 0.0177 - val_loss: 0.0188
Epoch 4/100
2/2 [=====] - 0s 43ms/step - loss: 0.0198 - val_loss: 0.0185
Epoch 5/100
2/2 [=====] - 0s 48ms/step - loss: 0.0195 - val_loss: 0.0187
Epoch 6/100
2/2 [=====] - 0s 51ms/step - loss: 0.0193 - val_loss: 0.0197
Epoch 7/100
2/2 [=====] - 0s 74ms/step - loss: 0.0240 - val_loss: 0.0214
Epoch 8/100
2/2 [=====] - 0s 102ms/step - loss: 0.0214 - val_loss: 0.0237
Epoch 9/100
2/2 [=====] - 0s 45ms/step - loss: 0.0226 - val_loss: 0.0256
Epoch 10/100
2/2 [=====] - 0s 60ms/step - loss: 0.0199 - val_loss: 0.0265
Epoch 11/100
2/2 [=====] - 0s 39ms/step - loss: 0.0199 - val_loss: 0.0274
Epoch 12/100
2/2 [=====] - 0s 41ms/step - loss: 0.0195 - val_loss: 0.0273
Epoch 13/100
2/2 [=====] - 0s 119ms/step - loss: 0.0218 - val_loss: 0.0265
Epoch 14/100
2/2 [=====] - 0s 40ms/step - loss: 0.0251 - val_loss: 0.0253
```

```
In [66]: 1 history_bilstm = fit_model(model_bilstm)
```

```
Epoch 1/100
2/2 [=====] - 0s 84ms/step - loss: 0.0172 - val_loss: 0.0175
Epoch 2/100
2/2 [=====] - 0s 37ms/step - loss: 0.0175 - val_loss: 0.0175
Epoch 3/100
2/2 [=====] - 0s 37ms/step - loss: 0.0172 - val_loss: 0.0182
Epoch 4/100
2/2 [=====] - 0s 38ms/step - loss: 0.0165 - val_loss: 0.0198
Epoch 5/100
2/2 [=====] - 0s 33ms/step - loss: 0.0161 - val_loss: 0.0219
Epoch 6/100
2/2 [=====] - 0s 31ms/step - loss: 0.0160 - val_loss: 0.0240
Epoch 7/100
2/2 [=====] - 0s 30ms/step - loss: 0.0161 - val_loss: 0.0253
Epoch 8/100
2/2 [=====] - 0s 31ms/step - loss: 0.0161 - val_loss: 0.0254
Epoch 9/100
2/2 [=====] - 0s 31ms/step - loss: 0.0160 - val_loss: 0.0244
Epoch 10/100
2/2 [=====] - 0s 31ms/step - loss: 0.0157 - val_loss: 0.0228
Epoch 11/100
2/2 [=====] - 0s 30ms/step - loss: 0.0155 - val_loss: 0.0211
```

```
In [58]: 1 def plot_loss (history, model_name):
2         plt.figure(figsize = (10, 6))
3         plt.plot(history.history['loss'])
4         plt.plot(history.history['val_loss'])
5         plt.title('Model Train vs Validation Loss for ' + model_name)
6         plt.ylabel('Loss')
7         plt.xlabel('epoch')
8         plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10        plot_loss (history_gru, 'GRU')
11        plot_loss (history_bilstm, 'Bidirectional LSTM')
```



In [62]:

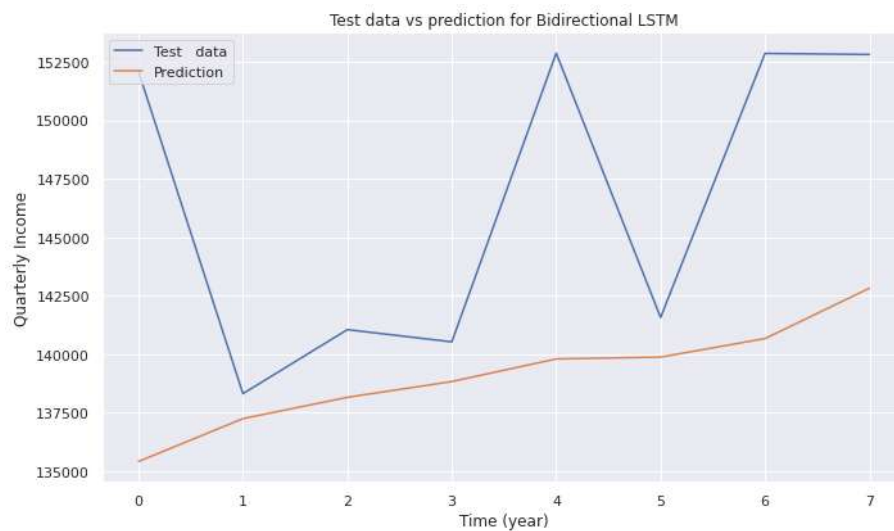

```

1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)

```

1/1 [=====] - 0s 39ms/step

1/1 [=====] - 0s 76ms/step



```
In [69]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11        evaluate_prediction(prediction_lstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')
```

```
In [78]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11        evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
  warnings.warn(
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
  warnings.warn(

```

1