

## Predicting the Quarterly Gross Profit for Amazon

```
In [10]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [11]: 1 # Reading the Data
2 df=pd.read_excel('Amazon Quarterly Gross Profit.xlsx')
3 df.head()
```

```
Out[11]:
```

	Date	Quarterly Profit
0	2009-03-31	1148
1	2009-06-30	1133
2	2009-09-30	1273
3	2009-12-31	1977
4	2010-03-31	1630

```
In [12]: 1 df.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55 entries, 0 to 54
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            55 non-null    datetime64[ns]
1   Quarterly Profit 55 non-null    int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 1008.0 bytes
```

```
In [13]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

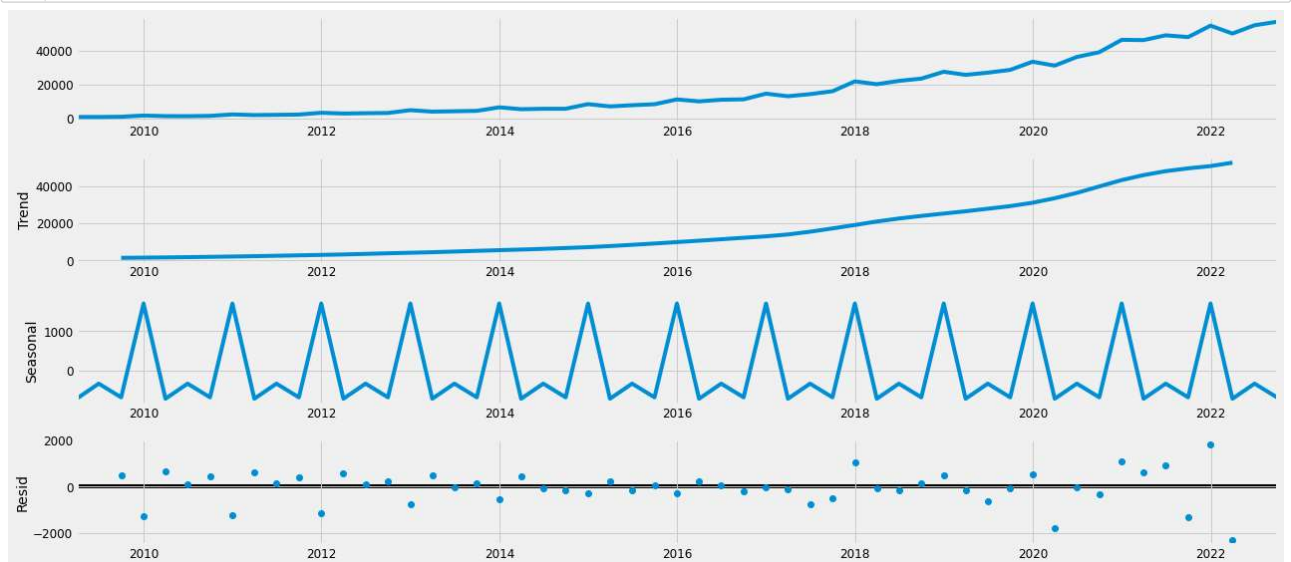
```
Out[13]:
```

	Quarterly Profit
Date	
2009-03-31	1148
2009-06-30	1133
2009-09-30	1273
2009-12-31	1977
2010-03-31	1630

```
In [14]: 1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Amazon Quarterly Gross Profit')
4 plt.show()
```



```
In [15]: 1 # Decomposition the data
2 from pylab import rcParams
3 rcParams['figure.figsize'] = 18, 8
4
5 decomposition = sm.tsa.seasonal_decompose(df, model = 'additive')
6 fig = decomposition.plot()
7 plt.show()
```



## SARIMA Model

```
In [17]: 1 df=pd.read_excel('Amazon Quarterly Gross Profit.xlsx')
2 df = df.set_index('Date')
```

```
In [18]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)  
SARIMAX: (0, 0, 1) x (0, 1, 0, 12)  
SARIMAX: (0, 1, 0) x (0, 1, 1, 12)  
SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [19]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_trend = False)
6             result = mod.fit()
7             print('SARIMA{0}{1}12 - AIC:{0}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)

```
In [20]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                   seasonal_order = (1, 1, 0, 12)
5                                   )
6 result = model.fit()
7 print(result.summary().tables[1])
```

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)  
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.  
self.\_init\_dates(dates, freq)

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          -0.4749        0.492       -0.965      0.334       -1.439        0.489
ma.L1           0.1957        0.587        0.333      0.739       -0.955        1.346
ar.S.L12         0.8051        0.203        3.973      0.000         0.408        1.202
sigma2          1.255e+06      2.32e+05        5.421      0.000      8.02e+05      1.71e+06
=====
```

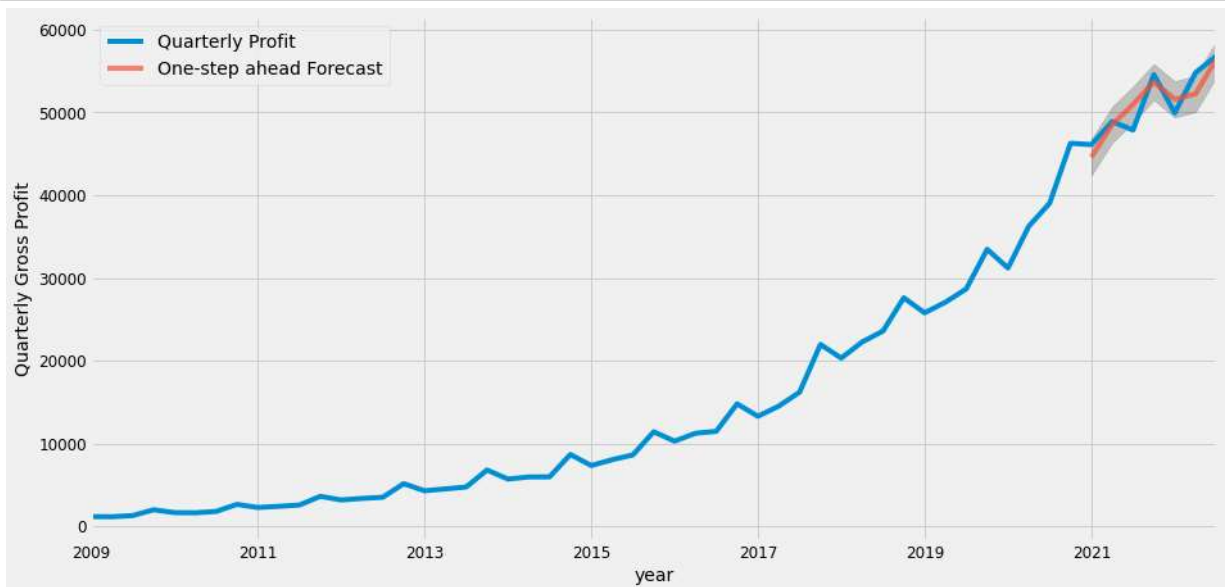
```
In [21]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-03-31'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

Out[21]:

	lower Quarterly Profit	upper Quarterly Profit
--	------------------------	------------------------

2021-03-31	42398.000144	46790.208160
2021-06-30	46333.307402	50725.515418
2021-09-30	48772.541951	53164.749967
2021-12-31	51484.935925	55877.143941
2022-03-31	49395.266486	53787.474502
2022-06-30	50054.645715	54446.853731
2022-09-30	54264.242912	58656.450928

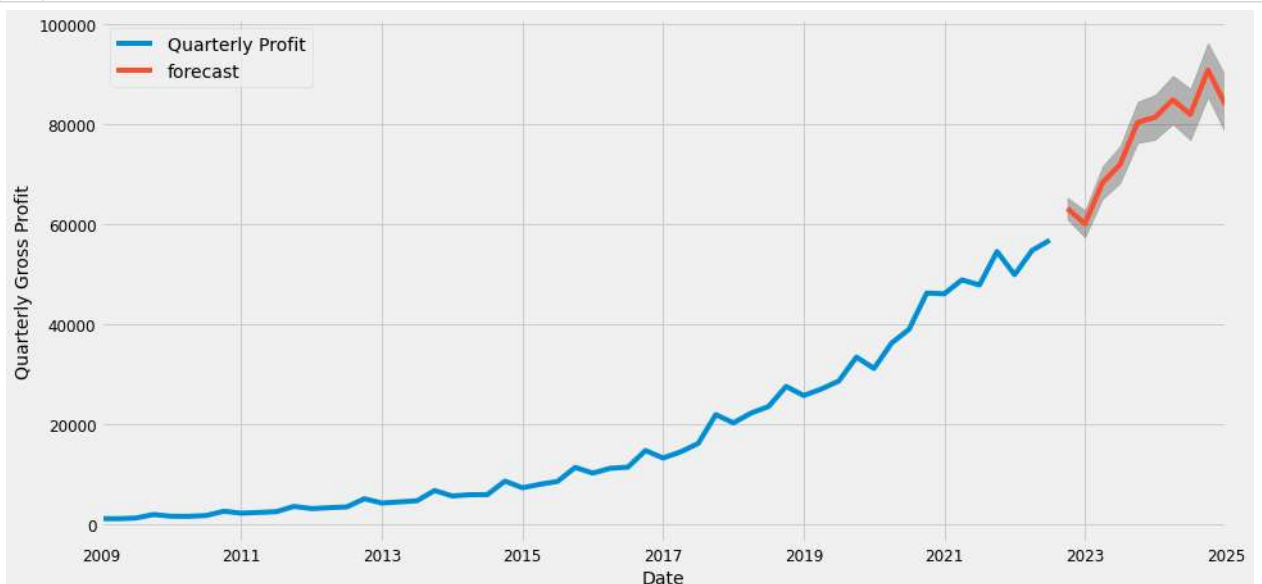
```
In [22]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Gross Profit')
7 plt.legend()
8 plt.show()
```



```
In [23]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-03-31:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 3168962.372852964  
The Root Mean Squared Error of our forecasts is 1780.1579628934517

```
In [24]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Gross Profit')
10
11 plt.legend()
12 plt.show()
13
```



## DNN MODEL

```
In [25]: 1 def convert2matrix(data_arr, look_back):
2         X, Y = [], []
3         for i in range(len(data_arr)-look_back):
4             d=i+look_back
5             X.append(data_arr[i:d,0])
6             Y.append(data_arr[d,0])
7         return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [26]: 1 df=pd.read_excel('Amazon Quarterly Gross Profit.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [27]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [28]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [29]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)
```

```
Epoch 1/500
12/12 [=====] - 1s 16ms/step - loss: 2350840.2500 - mse: 2350840.2500 - mae: 1033.5248 - val_loss: 54
679544.0000 - val_mse: 54679544.0000 - val_mae: 7304.3223
Epoch 2/500
12/12 [=====] - 0s 4ms/step - loss: 2116283.5000 - mse: 2116283.5000 - mae: 986.0873 - val_loss: 5292
0608.0000 - val_mse: 52920608.0000 - val_mae: 7196.0234
Epoch 3/500
12/12 [=====] - 0s 4ms/step - loss: 1979505.3750 - mse: 1979505.3750 - mae: 952.7255 - val_loss: 4929
4968.0000 - val_mse: 49294968.0000 - val_mae: 6949.4453
Epoch 4/500
12/12 [=====] - 0s 4ms/step - loss: 2070286.6250 - mse: 2070286.6250 - mae: 966.9977 - val_loss: 5102
0464.0000 - val_mse: 51020464.0000 - val_mae: 7077.1367
Epoch 5/500
12/12 [=====] - 0s 3ms/step - loss: 1921211.3750 - mse: 1921211.3750 - mae: 931.6425 - val_loss: 4565
1016.0000 - val_mse: 45651016.0000 - val_mae: 6691.1348
Epoch 6/500
12/12 [=====] - 0s 3ms/step - loss: 2239483.5000 - mse: 2239483.5000 - mae: 992.1248 - val_loss: 5280
2632.0000 - val_mse: 52802632.0000 - val_mae: 7206.4922
Epoch 7/500
12/12 [=====] - 0s 3ms/step - loss: 2012070.0000 - mse: 2012070.0000 - mae: 934.6333 - val_loss: 4037
1216.0000 - val_mse: 40371216.0000 - val_mae: 6151.0078
```


```
In [30]: 1 def model_loss(history):
2         plt.figure(figsize=(8,4))
3         plt.plot(history.history['loss'], label='Train Loss')
4         plt.plot(history.history['val_loss'], label='Test Loss')
5         plt.title('model loss')
6         plt.ylabel('loss')
7         plt.xlabel('epochs')
8         plt.legend(loc='upper right')
9         plt.show();
```

```
In [31]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))
```

```
Train Root Mean Squared Error(RMSE): 1158.29; Train Mean Absolute Error(MAE) : 734.11
[1341638.25, 1341638.25, 734.110778085938]
[38831080.0, 38831080.0, 6151.0078125]
Test Root Mean Squared Error(RMSE): 6231.46; Test Mean Absolute Error(MAE) : 6151.01
```

```
In [32]: 1 def prediction_plot(testY, test_predict):
2         len_prediction=[x for x in range(len(testY))]
3         plt.figure(figsize=(8,4))
4         plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5         plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6         plt.tight_layout()
7         sns.despine(top=True)
8         plt.subplots_adjust(left=0.07)
9         plt.ylabel('Ads Daily Spend', size=15)
10        plt.xlabel('Time step', size=15)
11        plt.legend(fontsize=15)
12        plt.show();
```

```
In [33]: 1 test_predict = model.predict(testX)
2         prediction_plot(testY, test_predict)
```



## GRU and BiLSTM Models

```
In [34]: 1 df=pd.read_excel('Amazon Quarterly Gross Profit.xlsx')
2         df.head()
```

```
Out[34]:
```

	Date	Quarterly Profit
0	2009-03-31	1148
1	2009-06-30	1133
2	2009-09-30	1273
3	2009-12-31	1977
4	2010-03-31	1630

```
In [35]: 1 df = df.set_index('Date')
2         df.head()
```

```
Out[35]:
```

	Quarterly Profit
Date	
2009-03-31	1148
2009-06-30	1133
2009-09-30	1273
2009-12-31	1977
2010-03-31	1630

```
In [36]: 1 # Split train data and test data
2         train_size = int(len(df)*0.8)
3
4         train_data = df.iloc[:train_size]
5         test_data = df.iloc[train_size:]
```

```
In [37]: 1 scaler = MinMaxScaler().fit(train_data)
2         train_scaled = scaler.transform(train_data)
3         test_scaled = scaler.transform(test_data)
```

```
In [38]: 1 # Create input dataset
2 def create_dataset (X, look_back = 1):
3     Xs, ys = [], []
4
5     for i in range(len(X)-look_back):
6         v = X[i:i+look_back]
7         Xs.append(v)
8         ys.append(X[i+look_back])
9
10    return np.array(Xs), np.array(ys)
11 LOOK_BACK = 4
12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
14 # Print data shape
15 print('X_train.shape: ', X_train.shape)
16 print('y_train.shape: ', y_train.shape)
17 print('X_test.shape: ', X_test.shape)
18 print('y_test.shape: ', y_test.shape)
```

```
X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (7, 4, 1)
y_test.shape: (7, 1)
```

```
In [39]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```

In [40]:

```
1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3                                                 patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5                         validation_split = 0.2,
6                         batch_size = 16, shuffle = False,
7                         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```



```
Epoch 1/100
2/2 [=====] - 3s 715ms/step - loss: 0.0467 - val_loss: 0.4321
Epoch 2/100
2/2 [=====] - 0s 35ms/step - loss: 0.0279 - val_loss: 0.3178
Epoch 3/100
2/2 [=====] - 0s 32ms/step - loss: 0.0155 - val_loss: 0.2220
Epoch 4/100
2/2 [=====] - 0s 31ms/step - loss: 0.0117 - val_loss: 0.1464
Epoch 5/100
2/2 [=====] - 0s 35ms/step - loss: 0.0097 - val_loss: 0.0932
Epoch 6/100
2/2 [=====] - 0s 32ms/step - loss: 0.0103 - val_loss: 0.0614
Epoch 7/100
2/2 [=====] - 0s 32ms/step - loss: 0.0134 - val_loss: 0.0457
Epoch 8/100
2/2 [=====] - 0s 37ms/step - loss: 0.0125 - val_loss: 0.0394
Epoch 9/100
2/2 [=====] - 0s 29ms/step - loss: 0.0114 - val_loss: 0.0380
Epoch 10/100
2/2 [=====] - 0s 28ms/step - loss: 0.0089 - val_loss: 0.0386
Epoch 11/100
2/2 [=====] - 0s 35ms/step - loss: 0.0077 - val_loss: 0.0388
Epoch 12/100
2/2 [=====] - 0s 37ms/step - loss: 0.0054 - val_loss: 0.0370
Epoch 13/100
2/2 [=====] - 0s 34ms/step - loss: 0.0048 - val_loss: 0.0325
Epoch 14/100
2/2 [=====] - 0s 35ms/step - loss: 0.0041 - val_loss: 0.0252
Epoch 15/100
2/2 [=====] - 0s 33ms/step - loss: 0.0038 - val_loss: 0.0161
Epoch 16/100
2/2 [=====] - 0s 37ms/step - loss: 0.0030 - val_loss: 0.0078
Epoch 17/100
2/2 [=====] - 0s 37ms/step - loss: 0.0028 - val_loss: 0.0030
Epoch 18/100
2/2 [=====] - 0s 33ms/step - loss: 0.0035 - val_loss: 0.0024
Epoch 19/100
2/2 [=====] - 0s 38ms/step - loss: 0.0013 - val_loss: 0.0049
Epoch 20/100
2/2 [=====] - 0s 35ms/step - loss: 0.0026 - val_loss: 0.0079
Epoch 21/100
2/2 [=====] - 0s 35ms/step - loss: 0.0017 - val_loss: 0.0098
Epoch 22/100
2/2 [=====] - 0s 35ms/step - loss: 0.0013 - val_loss: 0.0120
Epoch 23/100
2/2 [=====] - 0s 39ms/step - loss: 0.0021 - val_loss: 0.0138
Epoch 24/100
2/2 [=====] - 0s 55ms/step - loss: 0.0021 - val_loss: 0.0144
Epoch 25/100
2/2 [=====] - 0s 42ms/step - loss: 0.0025 - val_loss: 0.0129
Epoch 26/100
2/2 [=====] - 0s 42ms/step - loss: 0.0028 - val_loss: 0.0110
Epoch 27/100
2/2 [=====] - 0s 41ms/step - loss: 0.0025 - val_loss: 0.0110
Epoch 28/100
2/2 [=====] - 0s 44ms/step - loss: 0.0033 - val_loss: 0.0124
```



```
In [41]: 1 history_bilstm = fit_model(model_bilstm)
```



```
Epoch 1/100
2/2 [=====] - 6s 1s/step - loss: 0.0481 - val_loss: 0.4576
Epoch 2/100
2/2 [=====] - 0s 38ms/step - loss: 0.0347 - val_loss: 0.3571
Epoch 3/100
2/2 [=====] - 0s 40ms/step - loss: 0.0238 - val_loss: 0.2638
Epoch 4/100
2/2 [=====] - 0s 35ms/step - loss: 0.0154 - val_loss: 0.1807
Epoch 5/100
2/2 [=====] - 0s 33ms/step - loss: 0.0100 - val_loss: 0.1119
Epoch 6/100
2/2 [=====] - 0s 33ms/step - loss: 0.0080 - val_loss: 0.0622
Epoch 7/100
2/2 [=====] - 0s 33ms/step - loss: 0.0086 - val_loss: 0.0330
Epoch 8/100
2/2 [=====] - 0s 34ms/step - loss: 0.0099 - val_loss: 0.0200
Epoch 9/100
2/2 [=====] - 0s 34ms/step - loss: 0.0098 - val_loss: 0.0164
Epoch 10/100
2/2 [=====] - 0s 35ms/step - loss: 0.0083 - val_loss: 0.0178
Epoch 11/100
2/2 [=====] - 0s 35ms/step - loss: 0.0064 - val_loss: 0.0215
Epoch 12/100
2/2 [=====] - 0s 37ms/step - loss: 0.0049 - val_loss: 0.0247
Epoch 13/100
2/2 [=====] - 0s 35ms/step - loss: 0.0041 - val_loss: 0.0251
Epoch 14/100
2/2 [=====] - 0s 33ms/step - loss: 0.0037 - val_loss: 0.0214
Epoch 15/100
2/2 [=====] - 0s 35ms/step - loss: 0.0033 - val_loss: 0.0145
Epoch 16/100
2/2 [=====] - 0s 33ms/step - loss: 0.0026 - val_loss: 0.0069
Epoch 17/100
2/2 [=====] - 0s 34ms/step - loss: 0.0019 - val_loss: 0.0025
Epoch 18/100
2/2 [=====] - 0s 34ms/step - loss: 0.0013 - val_loss: 0.0044
Epoch 19/100
2/2 [=====] - 0s 33ms/step - loss: 0.0010 - val_loss: 0.0131
Epoch 20/100
2/2 [=====] - 0s 34ms/step - loss: 0.0012 - val_loss: 0.0240
Epoch 21/100
2/2 [=====] - 0s 35ms/step - loss: 0.0013 - val_loss: 0.0310
Epoch 22/100
2/2 [=====] - 0s 35ms/step - loss: 0.0012 - val_loss: 0.0318
Epoch 23/100
2/2 [=====] - 0s 37ms/step - loss: 0.0010 - val_loss: 0.0285
Epoch 24/100
2/2 [=====] - 0s 33ms/step - loss: 0.0010 - val_loss: 0.0248
Epoch 25/100
2/2 [=====] - 0s 32ms/step - loss: 0.0012 - val_loss: 0.0227
Epoch 26/100
2/2 [=====] - 0s 34ms/step - loss: 0.0012 - val_loss: 0.0225
Epoch 27/100
2/2 [=====] - 0s 37ms/step - loss: 0.0012 - val_loss: 0.0237
```

```
In [42]: 1 def plot_loss (history, model_name):
2         plt.figure(figsize = (10, 6))
3         plt.plot(history.history['loss'])
4         plt.plot(history.history['val_loss'])
5         plt.title('Model Train vs Validation Loss for ' + model_name)
6         plt.ylabel('Loss')
7         plt.xlabel('epoch')
8         plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10        plot_loss (history_gru, 'GRU')
11        plot_loss (history_bilstm, 'Bidirectional LSTM')
```



In [43]:

```
1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)
```



In [44]:

```
1 def evaluate_prediction(predictions, actual, model_name):
2     errors = predictions - actual
3     mse = np.square(errors).mean()
4     rmse = np.sqrt(mse)
5     mae = np.abs(errors).mean()
6     print(model_name + ':')
7     print('Mean Absolute Error: {:.4f}'.format(mae))
8     print('Root Mean Square Error: {:.4f}'.format(rmse))
9     print('')
10 evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11 evaluate_prediction(prediction_bilstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')
```

GRU:  
Mean Absolute Error: 10472.3912  
Root Mean Square Error: 11057.1284

Bidirectional LSTM:  
Mean Absolute Error: 18223.3047  
Root Mean Square Error: 18836.8948

```
In [45]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11        evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 0.3240

Root Mean Square Error: 0.3420

Bidirectional LSTM:

Mean Absolute Error: 0.5637

Root Mean Square Error: 0.5827

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names

warnings.warn(

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names

warnings.warn(