

Predicting the Quarterly Net Income for Amazon

```
In [8]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [9]: 1 # Reading the Data
2 df=pd.read_excel('Amazon Quarterly Net Income.xlsx')
3 df.head()
```

```
Out[9]:
```

	Date	Quarterly Net Income
0	2009-03-31	177
1	2009-06-30	142
2	2009-09-30	199
3	2009-12-31	384
4	2010-03-31	299

```
In [10]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55 entries, 0 to 54
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  55 non-null    datetime64[ns]
1   Quarterly Net Income  55 non-null    int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 1008.0 bytes
```

```
In [11]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

```
Out[11]:
```

	Quarterly Net Income
Date	
2009-03-31	177
2009-06-30	142
2009-09-30	199
2009-12-31	384
2010-03-31	299


In [12]:

```
1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Amazon Quarterly Net Income')
4 plt.show()
```



In [13]:

```
1 # Decomposition the data
2 from pylab import rcParams
3 rcParams['figure.figsize'] = 18, 8
4
5 decomposition = sm.tsa.seasonal_decompose(df, model = 'additive')
6 fig = decomposition.plot()
7 plt.show()
```




```
In [14]: 1 # Dividing the data into training and testing
2 # Plotting the data
3 import seaborn as sns
4 df['Date'] = df.index
5 train = df[df['Date'] < pd.to_datetime("2020-12", format='%Y-%m')]
6 train['train'] = train['Quarterly Net Income']
7 del train['Date']
8 del train['Quarterly Net Income']
9 test = df[df['Date'] >= pd.to_datetime("2020-12", format='%Y-%m')]
10 del test['Date']
11 test['test'] = test['Quarterly Net Income']
12 del test['Quarterly Net Income']
13 plt.plot(train, color = "black")
14 plt.plot(test, color = "red")
15 plt.title("Train/Test split for Quarterly Net Income")
16 plt.ylabel("Quarterly Net Income")
17 plt.xlabel('Date')
18 sns.set()
19 plt.show()
```

C:\Users\ravit\AppData\Local\Temp\ipykernel_9332\2151717043.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
train['train'] = train['Quarterly Net Income']
C:\Users\ravit\AppData\Local\Temp\ipykernel_9332\2151717043.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
test['test'] = test['Quarterly Net Income']



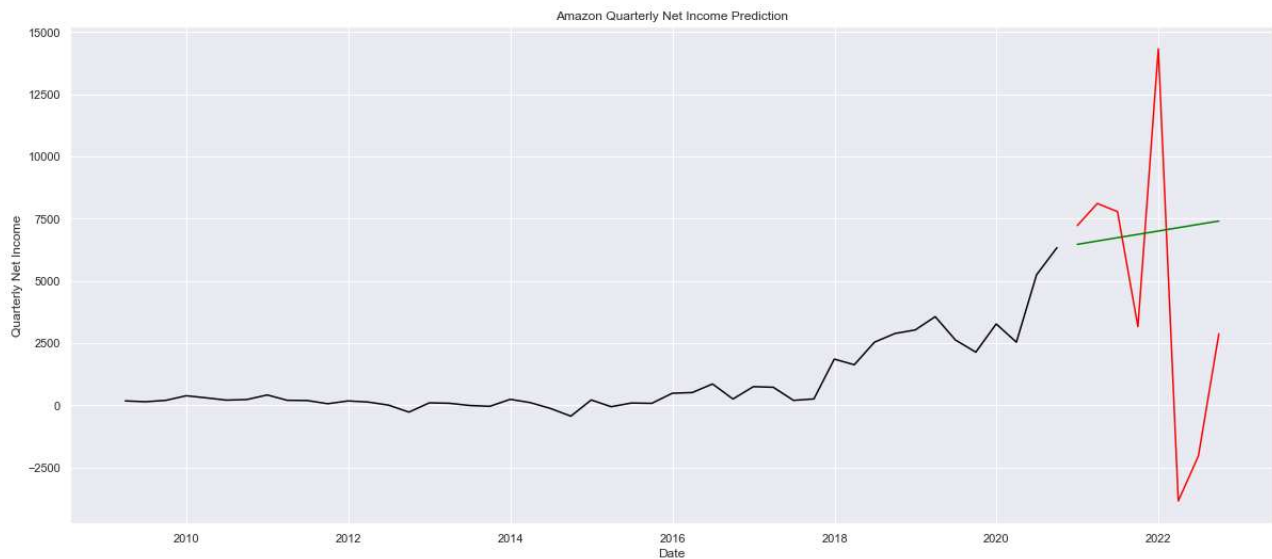
Arima Model

```
In [15]: 1 # Applying ARIMA Model
2 from pmdarima.arima import auto_arima
3 model = auto_arima(train, trace=True, error_action='ignore', suppress_warnings=True)
4 model.fit(train)
5 forecast = model.predict(n_periods=len(test))
6 forecast = pd.DataFrame(forecast, index = test.index, columns=['Prediction'])
```

Performing stepwise search to minimize aic
 ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.55 sec
 ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=723.843, Time=0.04 sec
 ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=725.025, Time=0.03 sec
 ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=725.129, Time=0.10 sec
 ARIMA(0,1,0)(0,0,0)[0] : AIC=724.038, Time=0.01 sec
 ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=726.967, Time=0.11 sec

Best model: ARIMA(0,1,0)(0,0,0)[0] intercept
 Total fit time: 0.864 seconds

```
In [16]: 1 # Plotting the prediction
2 plt.plot(train, color = "black")
3 plt.plot(test, color = "red")
4 plt.plot(forecast, color = "green")
5 plt.title("Amazon Quarterly Net Income Prediction")
6 plt.ylabel("Quarterly Net Income")
7 plt.xlabel('Date')
8 sns.set()
9 plt.show()
```



```
In [17]: 1 from math import sqrt
2 from sklearn.metrics import mean_squared_error
3 rms = sqrt(mean_squared_error(test,forecast))
4 print("RMSE: ", rms)
```

RMSE: 6111.103976680451

SARIMA Model

```
In [18]: 1 df=pd.read_excel('Amazon Quarterly Net Income.xlsx')
2 df = df.set_index('Date')
```

```
In [19]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)

SARIMAX: (0, 0, 1) x (0, 1, 0, 12)

SARIMAX: (0, 1, 0) x (0, 1, 1, 12)

SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [20]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_trend = False)
6             result = mod.fit()
7             print('SARIMA{x}{y}12 - AIC:{z}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

```
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
SARIMA(0, 0, 0)\(0, 0, 0) 12112 - AIC:1045.8693541174176
```

```
In [21]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                   seasonal_order = (1, 1, 0, 12)
5                                   )
6 result = model.fit()
7 print(result.summary().tables[1])
```

```
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: UserWarning: Non-stationary starting seasonal autoregressive Using zeros as starting parameters.
warn('Non-stationary starting seasonal autoregressive')
```


	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.2281	0.292	-0.781	0.435	-0.800	0.344
ma.L1	-0.5631	0.466	-1.208	0.227	-1.477	0.351
ar.S.L12	-0.0901	1.084	-0.083	0.934	-2.215	2.035
sigma2	9.071e+06	1.1e+06	8.245	0.000	6.91e+06	1.12e+07

```
In [22]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-03-31'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

Out[22]:

	lower Quarterly Net Income	upper Quarterly Net Income
2021-03-31	367.936860	12174.322671
2021-06-30	1752.300168	13558.685979
2021-09-30	2388.203534	14194.589345
2021-12-31	1438.612959	13244.998769
2022-03-31	2467.783946	14274.169757
2022-06-30	566.712621	12373.098431
2022-09-30	-4248.722171	7557.663640


```
In [23]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Net Income')
7 plt.legend()
8 plt.show()
```



```
In [25]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-03-31:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 43056747.27616465
The Root Mean Squared Error of our forecasts is 6561.764036916037

```
In [26]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Net Income')
10
11 plt.legend()
12 plt.show()
13
```



DNN MODEL

```
In [27]: 1 def convert2matrix(data_arr, look_back):
2         X, Y = [], []
3         for i in range(len(data_arr)-look_back):
4             d=i+look_back
5             X.append(data_arr[i:d,0])
6             Y.append(data_arr[d,0])
7         return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [28]: 1 df=pd.read_excel('Amazon Quarterly Net Income.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [29]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [30]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [31]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)
```

```
Epoch 1/500
12/12 [=====] - 2s 40ms/step - loss: 3621193.2500 - mse: 3621193.2500 - mae: 1064.9882 - val_loss: 14
530798.0000 - val_mse: 14530798.0000 - val_mae: 3766.7988
Epoch 2/500
12/12 [=====] - 0s 5ms/step - loss: 3066935.7500 - mse: 3066935.7500 - mae: 969.4407 - val_loss: 1625
3394.0000 - val_mse: 16253394.0000 - val_mae: 3868.6843
Epoch 3/500
12/12 [=====] - 0s 4ms/step - loss: 2573163.2500 - mse: 2573163.2500 - mae: 881.3557 - val_loss: 1857
1168.0000 - val_mse: 18571168.0000 - val_mae: 3979.4333
Epoch 4/500
12/12 [=====] - 0s 4ms/step - loss: 2135605.5000 - mse: 2135605.5000 - mae: 797.1470 - val_loss: 2146
7206.0000 - val_mse: 21467206.0000 - val_mae: 4111.5200
Epoch 5/500
12/12 [=====] - 0s 4ms/step - loss: 1764426.2500 - mse: 1764426.2500 - mae: 717.7985 - val_loss: 2477
3664.0000 - val_mse: 24773664.0000 - val_mae: 4254.0977
Epoch 6/500
12/12 [=====] - 0s 4ms/step - loss: 1446902.6250 - mse: 1446902.6250 - mae: 653.2609 - val_loss: 2848
5308.0000 - val_mse: 28485308.0000 - val_mae: 4410.8984
Epoch 7/500
12/12 [=====] - 0s 4ms/step - loss: 1165446.8750 - mse: 1165446.8750 - mae: 585.4536 - val_loss: 3274
1165446.8750 - val_mse: 1165446.8750 - val_mae: 585.4536
```


```
In [32]: 1 def model_loss(history):
2     plt.figure(figsize=(8,4))
3     plt.plot(history.history['loss'], label='Train Loss')
4     plt.plot(history.history['val_loss'], label='Test Loss')
5     plt.title('model loss')
6     plt.ylabel('loss')
7     plt.xlabel('epochs')
8     plt.legend(loc='upper right')
9     plt.show();
```

```
In [33]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))
```

```
Train Root Mean Squared Error(RMSE): 325.12; Train Mean Absolute Error(MAE) : 164.25
[105702.1796875, 105702.1796875, 164.24574279785156]
[336099744.0, 336099744.0, 13721.6064453125]
Test Root Mean Squared Error(RMSE): 18333.02; Test Mean Absolute Error(MAE) : 13721.61
```

```
In [34]: 1 def prediction_plot(testY, test_predict):
2         len_prediction=[x for x in range(len(testY))]
3         plt.figure(figsize=(8,4))
4         plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5         plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6         plt.tight_layout()
7         sns.despine(top=True)
8         plt.subplots_adjust(left=0.07)
9         plt.ylabel('Ads Daily Spend', size=15)
10        plt.xlabel('Time step', size=15)
11        plt.legend(fontsize=15)
12        plt.show();
```

```
In [35]: 1 test_predict = model.predict(testX)
2         prediction_plot(testY, test_predict)
```



GRU and BiLSTM Models

```
In [36]: 1 df=pd.read_excel('Amazon Quarterly Net Income.xlsx')
2         df.head()
```

```
Out[36]:
```

	Date	Quarterly Net Income
0	2009-03-31	177
1	2009-06-30	142
2	2009-09-30	199
3	2009-12-31	384
4	2010-03-31	299

```
In [37]: 1 df = df.set_index('Date')
2         df.head()
```

```
Out[37]:
```

	Date	Quarterly Net Income
	2009-03-31	177
	2009-06-30	142
	2009-09-30	199
	2009-12-31	384
	2010-03-31	299

```
In [38]: 1 # Split train data and test data
2         train_size = int(len(df)*0.8)
3
4         train_data = df.iloc[:train_size]
5         test_data = df.iloc[train_size:]
```

```
In [39]: 1 scaler = MinMaxScaler().fit(train_data)
2         train_scaled = scaler.transform(train_data)
3         test_scaled = scaler.transform(test_data)
```



```
In [40]: 1 # Create input dataset
2 def create_dataset (X, look_back = 1):
3     Xs, ys = [], []
4
5     for i in range(len(X)-look_back):
6         v = X[i:i+look_back]
7         Xs.append(v)
8         ys.append(X[i+look_back])
9
10    return np.array(Xs), np.array(ys)
11 LOOK_BACK = 4
12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
14 # Print data shape
15 print('X_train.shape: ', X_train.shape)
16 print('y_train.shape: ', y_train.shape)
17 print('X_test.shape: ', X_test.shape)
18 print('y_test.shape: ', y_test.shape)

X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (7, 4, 1)
y_test.shape: (7, 1)
```

```
In [41]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```

In [42]:

```
1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3                                                 patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5                         validation_split = 0.2,
6                         batch_size = 16, shuffle = False,
7                         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```

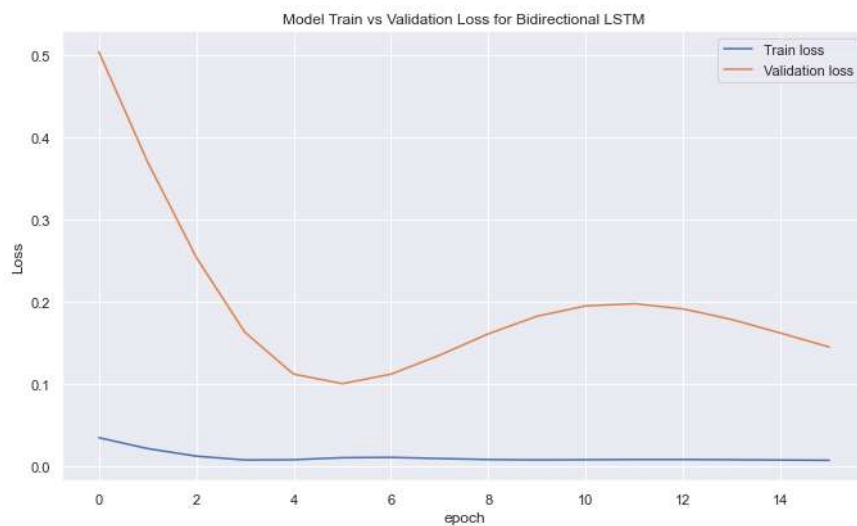
```
Epoch 1/100
2/2 [=====] - 4s 844ms/step - loss: 0.0311 - val_loss: 0.4045
Epoch 2/100
2/2 [=====] - 0s 31ms/step - loss: 0.0138 - val_loss: 0.2520
Epoch 3/100
2/2 [=====] - 0s 32ms/step - loss: 0.0084 - val_loss: 0.1504
Epoch 4/100
2/2 [=====] - 0s 29ms/step - loss: 0.0095 - val_loss: 0.1066
Epoch 5/100
2/2 [=====] - 0s 30ms/step - loss: 0.0147 - val_loss: 0.1066
Epoch 6/100
2/2 [=====] - 0s 39ms/step - loss: 0.0124 - val_loss: 0.1293
Epoch 7/100
2/2 [=====] - 0s 38ms/step - loss: 0.0100 - val_loss: 0.1597
Epoch 8/100
2/2 [=====] - 0s 41ms/step - loss: 0.0083 - val_loss: 0.1866
Epoch 9/100
2/2 [=====] - 0s 38ms/step - loss: 0.0092 - val_loss: 0.2030
Epoch 10/100
2/2 [=====] - 0s 49ms/step - loss: 0.0095 - val_loss: 0.2062
Epoch 11/100
2/2 [=====] - 0s 44ms/step - loss: 0.0087 - val_loss: 0.1977
Epoch 12/100
2/2 [=====] - 0s 40ms/step - loss: 0.0095 - val_loss: 0.1822
Epoch 13/100
2/2 [=====] - 0s 38ms/step - loss: 0.0087 - val_loss: 0.1630
Epoch 14/100
2/2 [=====] - 0s 44ms/step - loss: 0.0081 - val_loss: 0.1434
Epoch 15/100
2/2 [=====] - 0s 37ms/step - loss: 0.0089 - val_loss: 0.1272
```

In [43]:

```
1 history_bilstm = fit_model(model_bilstm)
```

```
Epoch 1/100
2/2 [=====] - 6s 2s/step - loss: 0.0349 - val_loss: 0.5041
Epoch 2/100
2/2 [=====] - 0s 39ms/step - loss: 0.0217 - val_loss: 0.3704
Epoch 3/100
2/2 [=====] - 0s 32ms/step - loss: 0.0125 - val_loss: 0.2542
Epoch 4/100
2/2 [=====] - 0s 30ms/step - loss: 0.0078 - val_loss: 0.1630
Epoch 5/100
2/2 [=====] - 0s 31ms/step - loss: 0.0081 - val_loss: 0.1121
Epoch 6/100
2/2 [=====] - 0s 29ms/step - loss: 0.0106 - val_loss: 0.1006
Epoch 7/100
2/2 [=====] - 0s 32ms/step - loss: 0.0110 - val_loss: 0.1120
Epoch 8/100
2/2 [=====] - 0s 33ms/step - loss: 0.0096 - val_loss: 0.1352
Epoch 9/100
2/2 [=====] - 0s 28ms/step - loss: 0.0083 - val_loss: 0.1611
Epoch 10/100
2/2 [=====] - 0s 37ms/step - loss: 0.0079 - val_loss: 0.1825
Epoch 11/100
2/2 [=====] - 0s 35ms/step - loss: 0.0081 - val_loss: 0.1951
Epoch 12/100
2/2 [=====] - 0s 34ms/step - loss: 0.0084 - val_loss: 0.1977
Epoch 13/100
2/2 [=====] - 0s 32ms/step - loss: 0.0084 - val_loss: 0.1914
Epoch 14/100
2/2 [=====] - 0s 34ms/step - loss: 0.0081 - val_loss: 0.1785
Epoch 15/100
2/2 [=====] - 0s 34ms/step - loss: 0.0077 - val_loss: 0.1620
Epoch 16/100
2/2 [=====] - 0s 33ms/step - loss: 0.0074 - val_loss: 0.1451
```

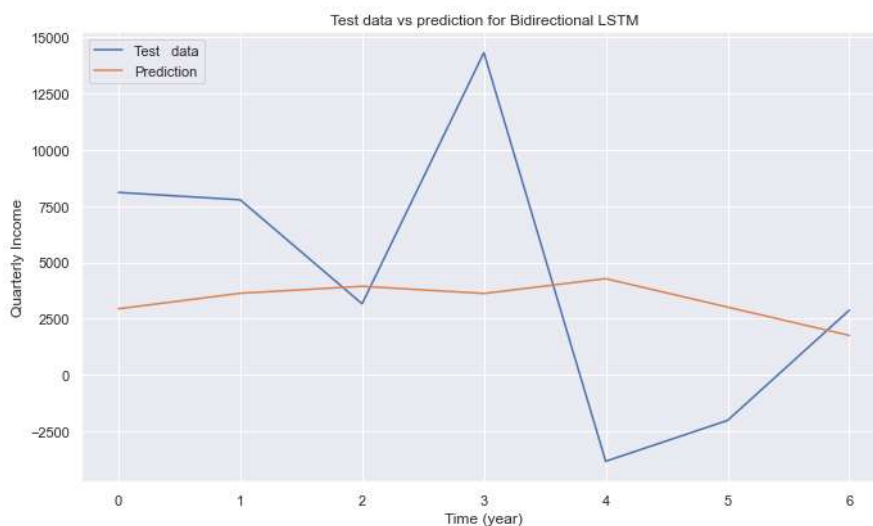
```
In [44]: 1 def plot_loss (history, model_name):
2         plt.figure(figsize = (10, 6))
3         plt.plot(history.history['loss'])
4         plt.plot(history.history['val_loss'])
5         plt.title('Model Train vs Validation Loss for ' + model_name)
6         plt.ylabel('Loss')
7         plt.xlabel('epoch')
8         plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10        plot_loss (history_gru, 'GRU')
11        plot_loss (history_bilstm, 'Bidirectional LSTM')
```



```

In [45]: 1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)

```



```

In [46]: 1 def evaluate_prediction(predictions, actual, model_name):
2     errors = predictions - actual
3     mse = np.square(errors).mean()
4     rmse = np.sqrt(mse)
5     mae = np.abs(errors).mean()
6     print(model_name + ':')
7     print('Mean Absolute Error: {:.4f}'.format(mae))
8     print('Root Mean Square Error: {:.4f}'.format(rmse))
9     print('')
10 evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11 evaluate_prediction(prediction_bilstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')

```

GRU:
Mean Absolute Error: 4998.6921
Root Mean Square Error: 5941.3557

Bidirectional LSTM:
Mean Absolute Error: 5012.3960
Root Mean Square Error: 5997.4695

```
In [47]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11        evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 1.2503

Root Mean Square Error: 1.4861

Bidirectional LSTM:

Mean Absolute Error: 1.2537

Root Mean Square Error: 1.5001

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

In []:

1