

Predicting the Quarterly Revenue for Amazon

```
In [1]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [2]: 1 # Reading the Data
2 df=pd.read_excel('Amazon Quarterly Revenue.xlsx')
3 df.head()
```

Out[2]:

	Date	Quarterly Revenue
0	2009-03-31	4889
1	2009-06-30	4651
2	2009-09-30	5449
3	2009-12-31	9520
4	2010-03-31	7131

```
In [3]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55 entries, 0 to 54
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  55 non-null    datetime64[ns]
1   Quarterly Revenue     55 non-null    int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 1008.0 bytes
```

```
In [4]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

Out[4]:

	Quarterly Revenue
Date	
2009-03-31	4889
2009-06-30	4651
2009-09-30	5449
2009-12-31	9520
2010-03-31	7131


In [5]:

```
1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Amazon Quarterly Revenue')
4 plt.show()
```



In [6]:

```
1 # Decomposition the data
2 from pylab import rcParams
3 rcParams['figure.figsize'] = 18, 8
4
5 decomposition = sm.tsa.seasonal_decompose(df, model = 'additive')
6 fig = decomposition.plot()
7 plt.show()
```



```
In [7]: 1 # Dividing the data into training and testing
2 # Plotting the data
3 import seaborn as sns
4 df['Date'] = df.index
5 train = df[df['Date'] < pd.to_datetime("2020-12", format='%Y-%m')]
6 train['train'] = train['Quarterly Revenue']
7 del train['Date']
8 del train['Quarterly Revenue']
9 test = df[df['Date'] >= pd.to_datetime("2020-12", format='%Y-%m')]
10 del test['Date']
11 test['test'] = test['Quarterly Revenue']
12 del test['Quarterly Revenue']
13 plt.plot(train, color = "black")
14 plt.plot(test, color = "red")
15 plt.title("Train/Test split for Quarterly Revenue")
16 plt.ylabel("Quarterly Revenue")
17 plt.xlabel('Date')
18 sns.set()
19 plt.show()
```

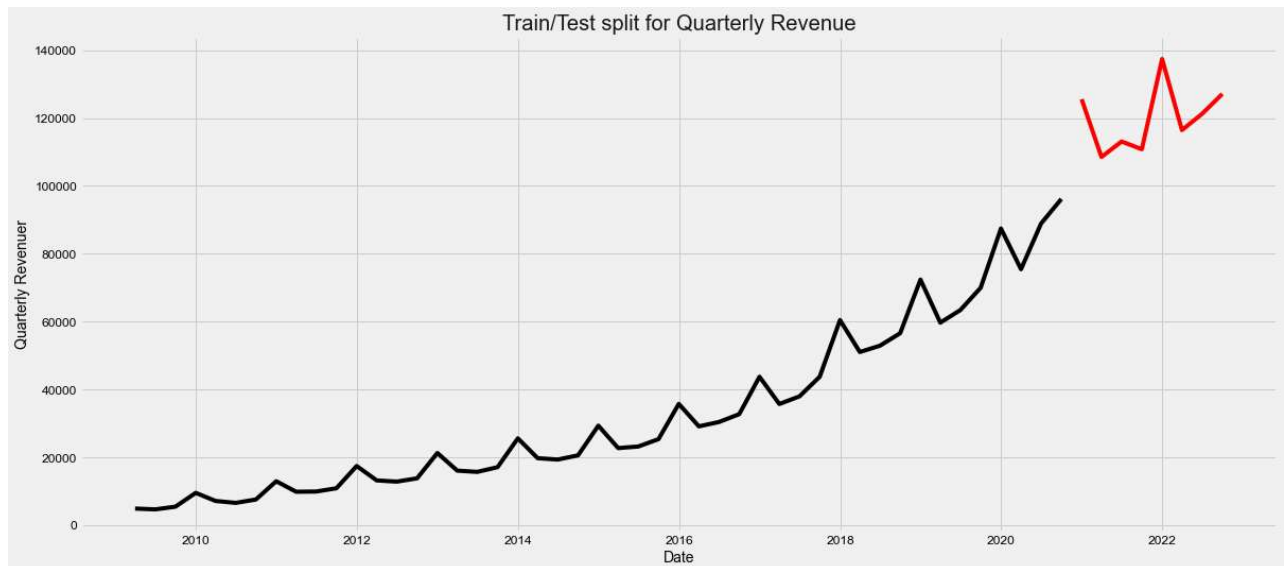
C:\Users\ravit\AppData\Local\Temp\ipykernel_23444\1521455262.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
train['train'] = train['Quarterly Revenue']
C:\Users\ravit\AppData\Local\Temp\ipykernel_23444\1521455262.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['test'] = test['Quarterly Revenue']
```



Arima Model

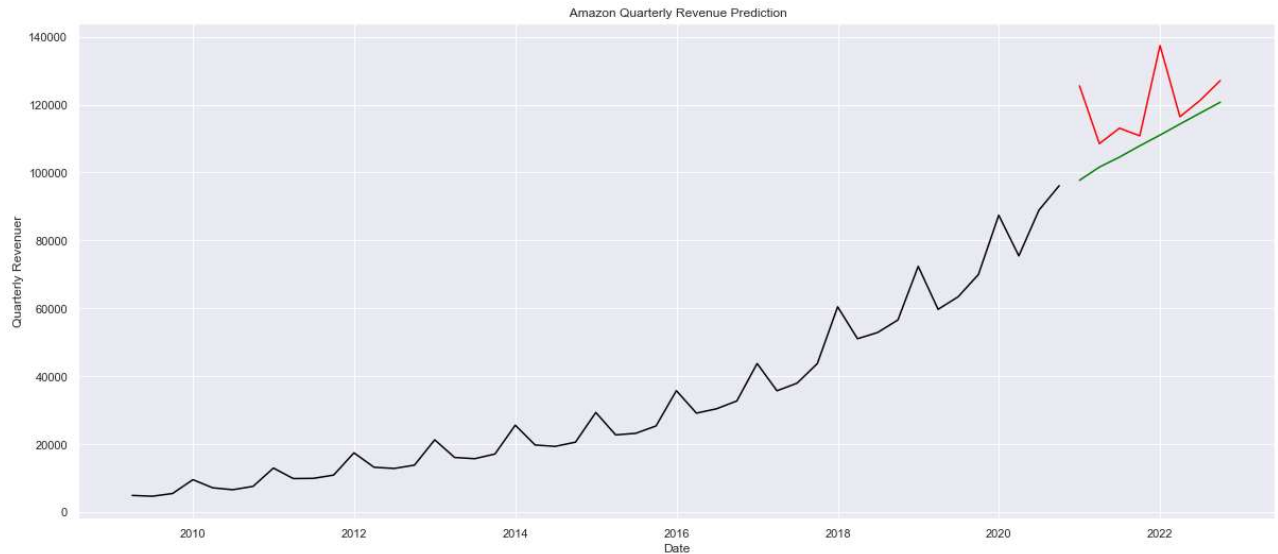
```
In [8]: 1 # Applying ARIMA Model
2 from pmdarima.arima import auto_arima
3 model = auto_arima(train, trace=True, error_action='ignore', suppress_warnings=True)
4 model.fit(train)
5 forecast = model.predict(n_periods=len(test))
6 forecast = pd.DataFrame(forecast, index = test.index, columns=['Prediction'])
```

Performing stepwise search to minimize aic

ARIMA(2,2,2)(0,0,0)[0]	: AIC=inf, Time=0.23 sec
ARIMA(0,2,0)(0,0,0)[0]	: AIC=971.654, Time=0.01 sec
ARIMA(1,2,0)(0,0,0)[0]	: AIC=952.220, Time=0.02 sec
ARIMA(0,2,1)(0,0,0)[0]	: AIC=inf, Time=0.05 sec
ARIMA(2,2,0)(0,0,0)[0]	: AIC=949.387, Time=0.02 sec
ARIMA(3,2,0)(0,0,0)[0]	: AIC=inf, Time=0.17 sec
ARIMA(2,2,1)(0,0,0)[0]	: AIC=inf, Time=0.12 sec
ARIMA(1,2,1)(0,0,0)[0]	: AIC=924.973, Time=0.04 sec
ARIMA(1,2,2)(0,0,0)[0]	: AIC=inf, Time=0.14 sec
ARIMA(0,2,2)(0,0,0)[0]	: AIC=inf, Time=0.08 sec
ARIMA(1,2,1)(0,0,0)[0] intercept	: AIC=inf, Time=0.16 sec

Best model: ARIMA(1,2,1)(0,0,0)[0]
Total fit time: 1.042 seconds

```
In [9]: 1 # Plotting the prediction
2 plt.plot(train, color = "black")
3 plt.plot(test, color = "red")
4 plt.plot(forecast, color = "green")
5 plt.title(" Amazon Quarterly Revenue Prediction")
6 plt.ylabel("Quarterly Revenuer")
7 plt.xlabel('Date')
8 sns.set()
9 plt.show()
```



```
In [10]: 1 from math import sqrt
2 from sklearn.metrics import mean_squared_error
3 rms = sqrt(mean_squared_error(test,forecast))
4 print("RMSE: ", rms)
```

RMSE: 14397.60532056588

SARIMA Model

```
In [11]: 1 df=pd.read_excel('Amazon Quarterly Revenue.xlsx')
2 df = df.set_index('Date')
3
```

```
In [12]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)

SARIMAX: (0, 0, 1) x (0, 1, 0, 12)

SARIMAX: (0, 1, 0) x (0, 1, 1, 12)

SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [13]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_trend = False)
6             result = mod.fit()
7             print('SARIMA{x}{y}12 - AIC:{z}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

SARIMA(0, 0, 0)x(0, 0, 0, 12)12 - AIC:1367.2792850832486

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)

```
In [14]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                   seasonal_order = (1, 1, 0, 12)
5                                   )
6 result = model.fit()
7 print(result.summary().tables[1])
```

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-DEC will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
warn('Non-stationary starting autoregressive parameters')
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: UserWarning: Non-stationary starting seasonal autoregressive Using zeros as starting parameters.
warn('Non-stationary starting seasonal autoregressive')

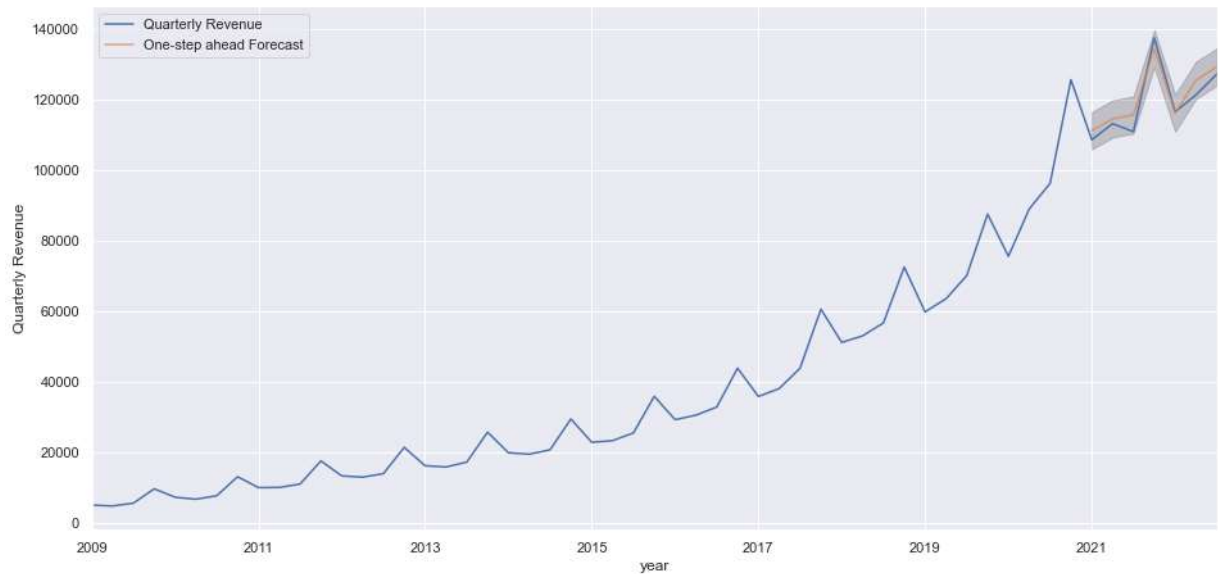
```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          -0.9427         0.140       -6.720      0.000       -1.218       -0.668
ma.L1           0.7850         0.263        2.983      0.003         0.269        1.301
ar.S.L12         0.8606         0.142        6.053      0.000         0.582        1.139
sigma2         7.27e+06      9.43e+09      7.71e+14      0.000      7.27e+06      7.27e+06
=====
```

```
In [16]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-03-31'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

Out[16]:

	lower Quarterly Revenue	upper Quarterly Revenue
2021-03-31	105815.138599	116384.229973
2021-06-30	109152.276180	119721.367405
2021-09-30	110279.879532	120848.970664
2021-12-31	129114.345905	139683.436980
2022-03-31	110812.233037	121381.324076
2022-06-30	120115.384149	130684.475167
2022-09-30	123874.142821	134443.233826


```
In [17]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Revenue')
7 plt.legend()
8 plt.show()
```



```
In [23]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-03-31:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 8840005.78536543
The Root Mean Squared Error of our forecasts is 2973.214722378024

```
In [24]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Revenue')
10
11 plt.legend()
12 plt.show()
13
```



DNN MODEL

```
In [25]: 1 def convert2matrix(data_arr, look_back):
2         X, Y = [], []
3         for i in range(len(data_arr)-look_back):
4             d=i+look_back
5             X.append(data_arr[i:d,0])
6             Y.append(data_arr[d,0])
7         return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [26]: 1 df=pd.read_excel('Amazon Quarterly Revenue.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [27]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [28]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [29]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)
```

```
Epoch 1/500
12/12 [=====] - 2s 22ms/step - loss: 1201522048.0000 - mse: 1201522048.0000 - mae: 27822.0586 - val_loss: 5053336064.0000 - val_mse: 5053336064.0000 - val_mae: 70970.8594
Epoch 2/500
12/12 [=====] - 0s 4ms/step - loss: 831125184.0000 - mse: 831125184.0000 - mae: 22757.4551 - val_loss: 2913479680.0000 - val_mse: 2913479680.0000 - val_mae: 53744.7383
Epoch 3/500
12/12 [=====] - 0s 3ms/step - loss: 542273984.0000 - mse: 542273984.0000 - mae: 18141.4531 - val_loss: 1351266816.0000 - val_mse: 1351266816.0000 - val_mae: 36260.3984
Epoch 4/500
12/12 [=====] - 0s 3ms/step - loss: 321139008.0000 - mse: 321139008.0000 - mae: 13588.3086 - val_loss: 407396864.0000 - val_mse: 407396864.0000 - val_mae: 19005.3984
Epoch 5/500
12/12 [=====] - 0s 3ms/step - loss: 165593744.0000 - mse: 165593744.0000 - mae: 9163.9678 - val_loss: 68919288.0000 - val_mse: 68919288.0000 - val_mae: 7771.8203
Epoch 6/500
12/12 [=====] - 0s 3ms/step - loss: 75645000.0000 - mse: 75645000.0000 - mae: 5967.2075 - val_loss: 200441792.0000 - val_mse: 200441792.0000 - val_mae: 11544.0547
Epoch 7/500
12/12 [=====] - 0s 3ms/step - loss: 32705016.0000 - mse: 32705016.0000 - mae: 4101.3500 - val_loss: 355905248.0000 - val_mse: 355905248.0000 - val_mae: 18217.0234
```


```
In [30]: 1 def model_loss(history):
2         plt.figure(figsize=(8,4))
3         plt.plot(history.history['loss'], label='Train Loss')
4         plt.plot(history.history['val_loss'], label='Test Loss')
5         plt.title('model loss')
6         plt.ylabel('loss')
7         plt.xlabel('epochs')
8         plt.legend(loc='upper right')
9         plt.show();
```

```
In [31]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))
```

```
Train Root Mean Squared Error(RMSE): 2863.85; Train Mean Absolute Error(MAE) : 1954.37
[8201623.0, 8201623.0, 1954.36669921875]
[355905248.0, 355905248.0, 18217.0234375]
Test Root Mean Squared Error(RMSE): 18865.45; Test Mean Absolute Error(MAE) : 18217.02
```

```
In [32]: 1 def prediction_plot(testY, test_predict):
2         len_prediction=[x for x in range(len(testY))]
3         plt.figure(figsize=(8,4))
4         plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5         plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6         plt.tight_layout()
7         sns.despine(top=True)
8         plt.subplots_adjust(left=0.07)
9         plt.ylabel('Ads Daily Spend', size=15)
10        plt.xlabel('Time step', size=15)
11        plt.legend(fontsize=15)
12        plt.show();
```

```
In [33]: 1 test_predict = model.predict(testX)
2         prediction_plot(testY, test_predict)
```



GRU and BiLSTM Models

```
In [34]: 1 df=pd.read_excel('Amazon Quarterly Revenue.xlsx')
2         df.head()
```

```
Out[34]:
```

	Date	Quarterly Revenue
0	2009-03-31	4889
1	2009-06-30	4651
2	2009-09-30	5449
3	2009-12-31	9520
4	2010-03-31	7131

```
In [35]: 1 df = df.set_index('Date')
2         df.head()
```

```
Out[35]:
```

	Date	Quarterly Revenue
	2009-03-31	4889
	2009-06-30	4651
	2009-09-30	5449
	2009-12-31	9520
	2010-03-31	7131

```
In [36]: 1 # Split train data and test data
2         train_size = int(len(df)*0.8)
3
4         train_data = df.iloc[:train_size]
5         test_data = df.iloc[train_size:]
```

```
In [37]: 1 scaler = MinMaxScaler().fit(train_data)
2         train_scaled = scaler.transform(train_data)
3         test_scaled = scaler.transform(test_data)
```



```
In [38]: 1 # Create input dataset
2 def create_dataset (X, look_back = 1):
3     Xs, ys = [], []
4
5     for i in range(len(X)-look_back):
6         v = X[i:i+look_back]
7         Xs.append(v)
8         ys.append(X[i+look_back])
9
10    return np.array(Xs), np.array(ys)
11 LOOK_BACK = 4
12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
14 # Print data shape
15 print('X_train.shape: ', X_train.shape)
16 print('y_train.shape: ', y_train.shape)
17 print('X_test.shape: ', X_test.shape)
18 print('y_test.shape: ', y_test.shape)
```

```
X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (7, 4, 1)
y_test.shape: (7, 1)
```

```
In [39]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```

In [40]:

```
1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3                                                 patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5                         validation_split = 0.2,
6                         batch_size = 16, shuffle = False,
7                         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```

```
Epoch 1/100
2/2 [=====] - 4s 639ms/step - loss: 0.0568 - val_loss: 0.3605
Epoch 2/100
2/2 [=====] - 0s 32ms/step - loss: 0.0325 - val_loss: 0.2366
Epoch 3/100
2/2 [=====] - 0s 29ms/step - loss: 0.0159 - val_loss: 0.1404
Epoch 4/100
2/2 [=====] - 0s 28ms/step - loss: 0.0121 - val_loss: 0.0744
Epoch 5/100
2/2 [=====] - 0s 29ms/step - loss: 0.0106 - val_loss: 0.0389
Epoch 6/100
2/2 [=====] - 0s 34ms/step - loss: 0.0139 - val_loss: 0.0255
Epoch 7/100
2/2 [=====] - 0s 29ms/step - loss: 0.0141 - val_loss: 0.0244
Epoch 8/100
2/2 [=====] - 0s 31ms/step - loss: 0.0132 - val_loss: 0.0302
Epoch 9/100
2/2 [=====] - 0s 31ms/step - loss: 0.0105 - val_loss: 0.0395
Epoch 10/100
2/2 [=====] - 0s 27ms/step - loss: 0.0111 - val_loss: 0.0490
Epoch 11/100
2/2 [=====] - 0s 31ms/step - loss: 0.0063 - val_loss: 0.0544
Epoch 12/100
2/2 [=====] - 0s 30ms/step - loss: 0.0063 - val_loss: 0.0540
Epoch 13/100
2/2 [=====] - 0s 33ms/step - loss: 0.0078 - val_loss: 0.0476
Epoch 14/100
2/2 [=====] - 0s 30ms/step - loss: 0.0074 - val_loss: 0.0377
Epoch 15/100
2/2 [=====] - 0s 30ms/step - loss: 0.0043 - val_loss: 0.0268
Epoch 16/100
2/2 [=====] - 0s 31ms/step - loss: 0.0040 - val_loss: 0.0169
Epoch 17/100
2/2 [=====] - 0s 29ms/step - loss: 0.0035 - val_loss: 0.0100
Epoch 18/100
2/2 [=====] - 0s 31ms/step - loss: 0.0047 - val_loss: 0.0069
Epoch 19/100
2/2 [=====] - 0s 32ms/step - loss: 0.0035 - val_loss: 0.0065
Epoch 20/100
2/2 [=====] - 0s 32ms/step - loss: 0.0044 - val_loss: 0.0072
Epoch 21/100
2/2 [=====] - 0s 32ms/step - loss: 0.0037 - val_loss: 0.0075
Epoch 22/100
2/2 [=====] - 0s 30ms/step - loss: 0.0022 - val_loss: 0.0073
Epoch 23/100
2/2 [=====] - 0s 31ms/step - loss: 0.0024 - val_loss: 0.0069
Epoch 24/100
2/2 [=====] - 0s 33ms/step - loss: 0.0028 - val_loss: 0.0066
Epoch 25/100
2/2 [=====] - 0s 30ms/step - loss: 0.0038 - val_loss: 0.0066
Epoch 26/100
2/2 [=====] - 0s 28ms/step - loss: 0.0028 - val_loss: 0.0068
Epoch 27/100
2/2 [=====] - 0s 31ms/step - loss: 0.0030 - val_loss: 0.0074
Epoch 28/100
2/2 [=====] - 0s 29ms/step - loss: 0.0045 - val_loss: 0.0088
Epoch 29/100
2/2 [=====] - 0s 29ms/step - loss: 0.0031 - val_loss: 0.0102
```



```
In [41]: 1 history_bilstm = fit_model(model_bilstm)
```

```
Epoch 1/100
2/2 [=====] - 5s 1s/step - loss: 0.0588 - val_loss: 0.3944
Epoch 2/100
2/2 [=====] - 0s 28ms/step - loss: 0.0404 - val_loss: 0.2915
Epoch 3/100
2/2 [=====] - 0s 27ms/step - loss: 0.0255 - val_loss: 0.1979
Epoch 4/100
2/2 [=====] - 0s 27ms/step - loss: 0.0144 - val_loss: 0.1182
Epoch 5/100
2/2 [=====] - 0s 27ms/step - loss: 0.0083 - val_loss: 0.0592
Epoch 6/100
2/2 [=====] - 0s 27ms/step - loss: 0.0076 - val_loss: 0.0257
Epoch 7/100
2/2 [=====] - 0s 29ms/step - loss: 0.0104 - val_loss: 0.0135
Epoch 8/100
2/2 [=====] - 0s 29ms/step - loss: 0.0121 - val_loss: 0.0117
Epoch 9/100
2/2 [=====] - 0s 28ms/step - loss: 0.0109 - val_loss: 0.0150
Epoch 10/100
2/2 [=====] - 0s 32ms/step - loss: 0.0083 - val_loss: 0.0223
Epoch 11/100
2/2 [=====] - 0s 32ms/step - loss: 0.0065 - val_loss: 0.0314
Epoch 12/100
2/2 [=====] - 0s 27ms/step - loss: 0.0058 - val_loss: 0.0383
Epoch 13/100
2/2 [=====] - 0s 30ms/step - loss: 0.0058 - val_loss: 0.0400
Epoch 14/100
2/2 [=====] - 0s 30ms/step - loss: 0.0057 - val_loss: 0.0358
Epoch 15/100
2/2 [=====] - 0s 31ms/step - loss: 0.0051 - val_loss: 0.0272
Epoch 16/100
2/2 [=====] - 0s 28ms/step - loss: 0.0043 - val_loss: 0.0172
Epoch 17/100
2/2 [=====] - 0s 29ms/step - loss: 0.0034 - val_loss: 0.0093
Epoch 18/100
2/2 [=====] - 0s 30ms/step - loss: 0.0028 - val_loss: 0.0057
Epoch 19/100
2/2 [=====] - 0s 33ms/step - loss: 0.0027 - val_loss: 0.0065
Epoch 20/100
2/2 [=====] - 0s 28ms/step - loss: 0.0028 - val_loss: 0.0089
Epoch 21/100
2/2 [=====] - 0s 30ms/step - loss: 0.0027 - val_loss: 0.0104
Epoch 22/100
2/2 [=====] - 0s 31ms/step - loss: 0.0024 - val_loss: 0.0101
Epoch 23/100
2/2 [=====] - 0s 33ms/step - loss: 0.0021 - val_loss: 0.0093
Epoch 24/100
2/2 [=====] - 0s 32ms/step - loss: 0.0021 - val_loss: 0.0089
Epoch 25/100
2/2 [=====] - 0s 32ms/step - loss: 0.0022 - val_loss: 0.0095
Epoch 26/100
2/2 [=====] - 0s 30ms/step - loss: 0.0022 - val_loss: 0.0115
Epoch 27/100
2/2 [=====] - 0s 31ms/step - loss: 0.0022 - val_loss: 0.0149
Epoch 28/100
2/2 [=====] - 0s 30ms/step - loss: 0.0022 - val_loss: 0.0188
```

```

In [42]: 1 def plot_loss (history, model_name):
2         plt.figure(figsize = (10, 6))
3         plt.plot(history.history['loss'])
4         plt.plot(history.history['val_loss'])
5         plt.title('Model Train vs Validation Loss for ' + model_name)
6         plt.ylabel('Loss')
7         plt.xlabel('epoch')
8         plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10        plot_loss (history_gru, 'GRU')
11        plot_loss (history_bilstm, 'Bidirectional LSTM')



```



```

In [43]: 1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)

```



```

In [44]: 1 def evaluate_prediction(predictions, actual, model_name):
2     errors = predictions - actual
3     mse = np.square(errors).mean()
4     rmse = np.sqrt(mse)
5     mae = np.abs(errors).mean()
6     print(model_name + ':')
7     print('Mean Absolute Error: {:.4f}'.format(mae))
8     print('Root Mean Square Error: {:.4f}'.format(rmse))
9     print('')
10 evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11 evaluate_prediction(prediction_bilstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')

```

GRU:
Mean Absolute Error: 23415.6786
Root Mean Square Error: 25093.8335

Bidirectional LSTM:
Mean Absolute Error: 44927.9888
Root Mean Square Error: 46282.9922

```
In [45]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10 evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11 evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 0.2828

Root Mean Square Error: 0.3031

Bidirectional LSTM:

Mean Absolute Error: 0.5427

Root Mean Square Error: 0.5591

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(