

Predicting the Quarterly Net Income for Walmart

```
In [1]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [2]: 1 # Reading the Data
2 df=pd.read_excel('Walmart Quarterly Net Income.xlsx')
3 df.head()
```

Out[2]:

	Date	Quarterly Net Income
0	2009-01-31	3772
1	2009-04-30	3022
2	2009-07-31	3472
3	2009-10-31	3144
4	2010-01-31	4732

```
In [3]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56 entries, 0 to 55
Data columns (total 2 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date                  56 non-null    datetime64[ns]
 1   Quarterly Net Income  56 non-null    int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 1.0 KB
```

```
In [4]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

Out[4]:

	Quarterly Net Income
Date	
2009-01-31	3772
2009-04-30	3022
2009-07-31	3472
2009-10-31	3144
2010-01-31	4732


In [5]:

```
1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Walmart Quarterly Revenue')
4 plt.show()
```



In [6]:

```
1 # Decomposition the data
2 from pylab import rcParams
3 rcParams['figure.figsize'] = 18, 8
4
5 decomposition = sm.tsa.seasonal_decompose(df, model = 'additive')
6 fig = decomposition.plot()
7 plt.show()
```




```
In [7]: 1 # Dividing the data into training and testing
2 # Plotting the data
3 import seaborn as sns
4 df['Date'] = df.index
5 train = df[df['Date'] < pd.to_datetime("2020-12", format='%Y-%m')]
6 train['train'] = train['Quarterly Net Income']
7 del train['Date']
8 del train['Quarterly Net Income']
9 test = df[df['Date'] >= pd.to_datetime("2020-12", format='%Y-%m')]
10 del test['Date']
11 test['test'] = test['Quarterly Net Income']
12 del test['Quarterly Net Income']
13 plt.plot(train, color = "black")
14 plt.plot(test, color = "red")
15 plt.title("Train/Test split for Quarterly Net Income")
16 plt.ylabel("Quarterly Net Income")
17 plt.xlabel('Date')
18 sns.set()
19 plt.show()
```

C:\Users\ravit\AppData\Local\Temp\ipykernel_20844\2151717043.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
train['train'] = train['Quarterly Net Income']

C:\Users\ravit\AppData\Local\Temp\ipykernel_20844\2151717043.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
test['test'] = test['Quarterly Net Income']



Arima Model

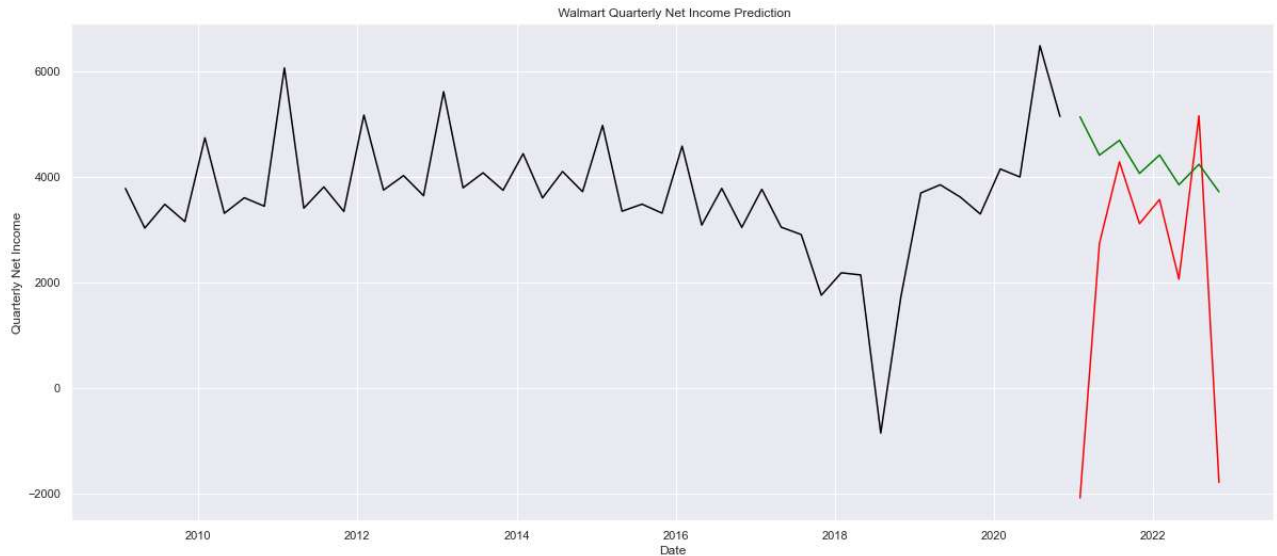
```
In [8]: 1 # Applying ARIMA Model
2 from pmdarima.arima import auto_arima
3 model = auto_arima(train, trace=True, error_action='ignore', suppress_warnings=True)
4 model.fit(train)
5 forecast = model.predict(n_periods=len(test))
6 forecast = pd.DataFrame(forecast, index = test.index, columns=['Prediction'])
```

Performing stepwise search to minimize aic

ARIMA(2,0,2)(0,0,0)[0] intercept	: AIC=806.808, Time=0.44 sec
ARIMA(0,0,0)(0,0,0)[0] intercept	: AIC=816.188, Time=0.02 sec
ARIMA(1,0,0)(0,0,0)[0] intercept	: AIC=810.220, Time=0.02 sec
ARIMA(0,0,1)(0,0,0)[0] intercept	: AIC=813.280, Time=0.09 sec
ARIMA(0,0,0)(0,0,0)[0]	: AIC=930.029, Time=0.01 sec
ARIMA(1,0,2)(0,0,0)[0] intercept	: AIC=813.033, Time=0.06 sec
ARIMA(2,0,1)(0,0,0)[0] intercept	: AIC=807.215, Time=0.28 sec
ARIMA(3,0,2)(0,0,0)[0] intercept	: AIC=808.238, Time=0.37 sec
ARIMA(2,0,3)(0,0,0)[0] intercept	: AIC=807.490, Time=0.28 sec
ARIMA(1,0,1)(0,0,0)[0] intercept	: AIC=810.846, Time=0.04 sec
ARIMA(1,0,3)(0,0,0)[0] intercept	: AIC=814.300, Time=0.10 sec
ARIMA(3,0,1)(0,0,0)[0] intercept	: AIC=808.793, Time=0.39 sec
ARIMA(3,0,3)(0,0,0)[0] intercept	: AIC=808.855, Time=0.40 sec
ARIMA(2,0,2)(0,0,0)[0]	: AIC=812.069, Time=0.21 sec

Best model: ARIMA(2,0,2)(0,0,0)[0] intercept
Total fit time: 2.739 seconds

```
In [9]: 1 # Plotting the prediction
2 plt.plot(train, color = "black")
3 plt.plot(test, color = "red")
4 plt.plot(forecast, color = "green")
5 plt.title(" Walmart Quarterly Net Income Prediction")
6 plt.ylabel("Quarterly Net Income")
7 plt.xlabel('Date')
8 sns.set()
9 plt.show()
```



```
In [10]: 1 from math import sqrt
2 from sklearn.metrics import mean_squared_error
3 rms = sqrt(mean_squared_error(test,forecast))
4 print("RMSE: ", rms)
```

RMSE: 3373.786162000279

SARIMA Model

```
In [11]: 1 df=pd.read_excel('Walmart Quarterly Net Income.xlsx')
2 df = df.set_index('Date')
```

```
In [12]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)
SARIMAX: (0, 0, 1) x (0, 1, 0, 12)
SARIMAX: (0, 1, 0) x (0, 1, 1, 12)
SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [13]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_invertibility = False)
6             result = mod.fit()
7             print('SARIMA{x}{y}12 - AIC:{z}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

```
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
```

```
In [14]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                     seasonal_order = (1, 1, 0, 12)
5                                     )
6 result = model.fit()
7 print(result.summary().tables[1])
```

```
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.
self._init_dates(dates, freq)
R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
warn('Non-invertible starting MA parameters found.')
```


```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.0253      0.307        0.082      0.934      -0.575      0.626
ma.L1         -0.7526      0.141       -5.332      0.000      -1.029     -0.476
ar.S.L12       -0.5014      0.218       -2.300      0.021      -0.929     -0.074
sigma2        3.299e+06    6.54e+05      5.045      0.000    2.02e+06    4.58e+06
=====
```

```
In [15]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-04-30'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

Out[15]:

	lower Quarterly Net Income	upper Quarterly Net Income
2021-04-30	-1654.421889	5465.365348
2021-07-31	-2733.527648	4386.259526
2021-10-31	-608.785230	6511.001910
2022-01-31	991.979169	8111.766289
2022-04-30	45.244418	7165.031526
2022-07-31	-127.355867	6992.431234
2022-10-31	-161.408001	6958.379097


```
In [16]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Net Income')
7 plt.legend()
8 plt.show()
```



```
In [17]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-04-30:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 6562726.660534659
The Root Mean Squared Error of our forecasts is 2561.7819307143727

```
In [18]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Net Income')
10
11 plt.legend()
12 plt.show()
13
```



DNN MODEL

```
In [19]: 1 def convert2matrix(data_arr, look_back):
2         X, Y = [], []
3         for i in range(len(data_arr)-look_back):
4             d=i+look_back
5             X.append(data_arr[i:d,0])
6             Y.append(data_arr[d,0])
7         return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [20]: 1 df=pd.read_excel('Walmart Quarterly Net Income.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [21]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [22]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [23]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)
```

```
Epoch 1/500
12/12 [=====] - 1s 21ms/step - loss: 10397825.0000 - mse: 10397825.0000 - mae: 3072.6726 - val_loss:
7986049.5000 - val_mse: 7986049.5000 - val_mae: 2522.6152
Epoch 2/500
12/12 [=====] - 0s 5ms/step - loss: 6700009.0000 - mse: 6700009.0000 - mae: 2413.0437 - val_loss: 848
9033.0000 - val_mse: 8489033.0000 - val_mae: 2491.5454
Epoch 3/500
12/12 [=====] - 0s 4ms/step - loss: 5305328.5000 - mse: 5305328.5000 - mae: 2086.3992 - val_loss: 853
7176.0000 - val_mse: 8537176.0000 - val_mae: 2424.1121
Epoch 4/500
12/12 [=====] - 0s 4ms/step - loss: 4452226.5000 - mse: 4452226.5000 - mae: 1854.4752 - val_loss: 870
2546.0000 - val_mse: 8702546.0000 - val_mae: 2435.6445
Epoch 5/500
12/12 [=====] - 0s 5ms/step - loss: 3831167.2500 - mse: 3831167.2500 - mae: 1662.9531 - val_loss: 902
3541.0000 - val_mse: 9023541.0000 - val_mae: 2541.0049
Epoch 6/500
12/12 [=====] - 0s 5ms/step - loss: 3359253.7500 - mse: 3359253.7500 - mae: 1490.4729 - val_loss: 952
5285.0000 - val_mse: 9525285.0000 - val_mae: 2646.3252
Epoch 7/500
12/12 [=====] - 0s 5ms/step - loss: 3007031.5000 - mse: 3007031.5000 - mae: 1331.0513 - val_loss: 1000
15373329.0 - val_mse: 15373329.0 - val_mae: 3059.4680
```

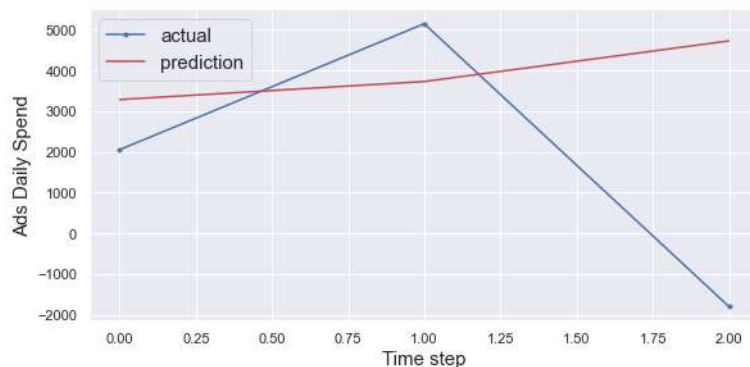
```
In [24]: 1 def model_loss(history):
2     plt.figure(figsize=(8,4))
3     plt.plot(history.history['loss'], label='Train Loss')
4     plt.plot(history.history['val_loss'], label='Test Loss')
5     plt.title('model loss')
6     plt.ylabel('loss')
7     plt.xlabel('epochs')
8     plt.legend(loc='upper right')
9     plt.show();
```

```
In [25]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))
```

```
Train Root Mean Squared Error(RMSE): 919.19; Train Mean Absolute Error(MAE) : 580.17
[844912.1875, 844912.1875, 580.17041015625]
[15373329.0, 15373329.0, 3059.468017578125]
Test Root Mean Squared Error(RMSE): 3920.88; Test Mean Absolute Error(MAE) : 3059.47
```

```
In [26]: 1 def prediction_plot(testY, test_predict):
2         len_prediction=[x for x in range(len(testY))]
3         plt.figure(figsize=(8,4))
4         plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5         plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6         plt.tight_layout()
7         sns.despine(top=True)
8         plt.subplots_adjust(left=0.07)
9         plt.ylabel('Ads Daily Spend', size=15)
10        plt.xlabel('Time step', size=15)
11        plt.legend(fontsize=15)
12        plt.show();
```

```
In [27]: 1 test_predict = model.predict(testX)
2         prediction_plot(testY, test_predict)
```



GRU and BiLSTM Models

```
In [28]: 1 df=pd.read_excel('Walmart Quarterly Net Income.xlsx')
2         df.head()
```

```
Out[28]:
```

	Date	Quarterly Net Income
0	2009-01-31	3772
1	2009-04-30	3022
2	2009-07-31	3472
3	2009-10-31	3144
4	2010-01-31	4732

```
In [29]: 1 df = df.set_index('Date')
2         df.head()
```

```
Out[29]:
```

	Date	Quarterly Net Income
	2009-01-31	3772
	2009-04-30	3022
	2009-07-31	3472
	2009-10-31	3144
	2010-01-31	4732

```
In [30]: 1 # Split train data and test data
2         train_size = int(len(df)*0.8)
3
4         train_data = df.iloc[:train_size]
5         test_data = df.iloc[train_size:]
```

```
In [31]: 1 scaler = MinMaxScaler().fit(train_data)
2         train_scaled = scaler.transform(train_data)
3         test_scaled = scaler.transform(test_data)
```



```
In [32]: 1 # Create input dataset
2 def create_dataset (X, look_back = 1):
3     Xs, ys = [], []
4
5     for i in range(len(X)-look_back):
6         v = X[i:i+look_back]
7         Xs.append(v)
8         ys.append(X[i+look_back])
9
10    return np.array(Xs), np.array(ys)
11 LOOK_BACK = 4
12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
14 # Print data shape
15 print('X_train.shape: ', X_train.shape)
16 print('y_train.shape: ', y_train.shape)
17 print('X_test.shape: ', X_test.shape)
18 print('y_test.shape: ', y_test.shape)
```

```
X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (8, 4, 1)
y_test.shape: (8, 1)
```

```
In [33]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```

In [34]:



```
1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3                                                 patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5                         validation_split = 0.2,
6                         batch_size = 16, shuffle = False,
7                         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```

```
Epoch 1/100
2/2 [=====] - 3s 606ms/step - loss: 0.3822 - val_loss: 0.1531
Epoch 2/100
2/2 [=====] - 0s 26ms/step - loss: 0.2012 - val_loss: 0.0856
Epoch 3/100
2/2 [=====] - 0s 39ms/step - loss: 0.0839 - val_loss: 0.0544
Epoch 4/100
2/2 [=====] - 0s 37ms/step - loss: 0.0290 - val_loss: 0.0613
Epoch 5/100
2/2 [=====] - 0s 29ms/step - loss: 0.0308 - val_loss: 0.0841
Epoch 6/100
2/2 [=====] - 0s 30ms/step - loss: 0.0540 - val_loss: 0.0916
Epoch 7/100
2/2 [=====] - 0s 32ms/step - loss: 0.0573 - val_loss: 0.0821
Epoch 8/100
2/2 [=====] - 0s 32ms/step - loss: 0.0433 - val_loss: 0.0669
Epoch 9/100
2/2 [=====] - 0s 29ms/step - loss: 0.0237 - val_loss: 0.0566
Epoch 10/100
2/2 [=====] - 0s 31ms/step - loss: 0.0199 - val_loss: 0.0526
Epoch 11/100
2/2 [=====] - 0s 30ms/step - loss: 0.0151 - val_loss: 0.0526
Epoch 12/100
2/2 [=====] - 0s 31ms/step - loss: 0.0209 - val_loss: 0.0540
Epoch 13/100
2/2 [=====] - 0s 28ms/step - loss: 0.0273 - val_loss: 0.0545
Epoch 14/100
2/2 [=====] - 0s 29ms/step - loss: 0.0256 - val_loss: 0.0538
Epoch 15/100
2/2 [=====] - 0s 30ms/step - loss: 0.0297 - val_loss: 0.0526
Epoch 16/100
2/2 [=====] - 0s 29ms/step - loss: 0.0165 - val_loss: 0.0521
Epoch 17/100
2/2 [=====] - 0s 29ms/step - loss: 0.0167 - val_loss: 0.0527
Epoch 18/100
2/2 [=====] - 0s 30ms/step - loss: 0.0165 - val_loss: 0.0544
Epoch 19/100
2/2 [=====] - 0s 30ms/step - loss: 0.0188 - val_loss: 0.0559
Epoch 20/100
2/2 [=====] - 0s 30ms/step - loss: 0.0198 - val_loss: 0.0565
Epoch 21/100
2/2 [=====] - 0s 31ms/step - loss: 0.0237 - val_loss: 0.0563
Epoch 22/100
2/2 [=====] - 0s 29ms/step - loss: 0.0207 - val_loss: 0.0556
Epoch 23/100
2/2 [=====] - 0s 29ms/step - loss: 0.0147 - val_loss: 0.0542
Epoch 24/100
2/2 [=====] - 0s 29ms/step - loss: 0.0170 - val_loss: 0.0531
Epoch 25/100
2/2 [=====] - 0s 30ms/step - loss: 0.0161 - val_loss: 0.0525
Epoch 26/100
2/2 [=====] - 0s 38ms/step - loss: 0.0141 - val_loss: 0.0522
```

```
In [35]: 1 history_bilstm = fit_model(model_bilstm)
```


```
Epoch 1/100
2/2 [=====] - 5s 1s/step - loss: 0.4587 - val_loss: 0.2108
Epoch 2/100
2/2 [=====] - 0s 33ms/step - loss: 0.3189 - val_loss: 0.1503
Epoch 3/100
2/2 [=====] - 0s 33ms/step - loss: 0.2007 - val_loss: 0.1008
Epoch 4/100
2/2 [=====] - 0s 34ms/step - loss: 0.1039 - val_loss: 0.0658
Epoch 5/100
2/2 [=====] - 0s 32ms/step - loss: 0.0374 - val_loss: 0.0535
Epoch 6/100
2/2 [=====] - 0s 32ms/step - loss: 0.0169 - val_loss: 0.0676
Epoch 7/100
2/2 [=====] - 0s 31ms/step - loss: 0.0407 - val_loss: 0.0840
Epoch 8/100
2/2 [=====] - 0s 31ms/step - loss: 0.0583 - val_loss: 0.0802
Epoch 9/100
2/2 [=====] - 0s 35ms/step - loss: 0.0450 - val_loss: 0.0669
Epoch 10/100
2/2 [=====] - 0s 30ms/step - loss: 0.0238 - val_loss: 0.0570
Epoch 11/100
2/2 [=====] - 0s 31ms/step - loss: 0.0129 - val_loss: 0.0537
Epoch 12/100
2/2 [=====] - 0s 30ms/step - loss: 0.0129 - val_loss: 0.0542
Epoch 13/100
2/2 [=====] - 0s 32ms/step - loss: 0.0177 - val_loss: 0.0556
Epoch 14/100
2/2 [=====] - 0s 29ms/step - loss: 0.0217 - val_loss: 0.0561
Epoch 15/100
2/2 [=====] - 0s 31ms/step - loss: 0.0228 - val_loss: 0.0555
```

```
In [36]: 1 def plot_loss (history, model_name):
2     plt.figure(figsize = (10, 6))
3     plt.plot(history.history['loss'])
4     plt.plot(history.history['val_loss'])
5     plt.title('Model Train vs Validation Loss for ' + model_name)
6     plt.ylabel('Loss')
7     plt.xlabel('epoch')
8     plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10 plot_loss (history_gru, 'GRU')
11 plot_loss (history_bilstm, 'Bidirectional LSTM')
```



In [37]:

```
1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)
```



In [38]:

```
1 def evaluate_prediction(predictions, actual, model_name):
2     errors = predictions - actual
3     mse = np.square(errors).mean()
4     rmse = np.sqrt(mse)
5     mae = np.abs(errors).mean()
6     print(model_name + ':')
7     print('Mean Absolute Error: {:.4f}'.format(mae))
8     print('Root Mean Square Error: {:.4f}'.format(rmse))
9     print('')
10 evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11 evaluate_prediction(prediction_bilstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 2311.8309

Root Mean Square Error: 3120.6002

Bidirectional LSTM:

Mean Absolute Error: 2268.0282

Root Mean Square Error: 2947.6046

```
In [39]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11        evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 0.3342

Root Mean Square Error: 0.4511

Bidirectional LSTM:

Mean Absolute Error: 0.3279

Root Mean Square Error: 0.4261

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

In []:

1