

Predicting the Quarterly Revenue for Walmart

```
In [24]: 1 # Importing Packages
2 import itertools
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 import matplotlib
8 import itertools
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import statsmodels.api as sm
13 import matplotlib
14 import sklearn.preprocessing
15 from sklearn.metrics import r2_score
16 import keras
17
18 from keras.layers import Dense,Dropout,SimpleRNN,GRU, Bidirectional,LSTM
19 from tensorflow.keras.optimizers import SGD
20 from keras.models import Sequential
21 from sklearn.preprocessing import MinMaxScaler, StandardScaler
22
23 plt.style.use('fivethirtyeight')
24 matplotlib.rcParams['axes.labelsize'] = 14
25 matplotlib.rcParams['xtick.labelsize'] = 12
26 matplotlib.rcParams['ytick.labelsize'] = 12
27 matplotlib.rcParams['text.color'] = 'k'
```

```
In [25]: 1 # Reading the Data
2 df=pd.read_excel('Walmart Quarterly Revenue.xlsx')
3 df.head()
```

```
Out[25]:
```

	Date	Quarterly Revenue
0	2009-01-31	108627
1	2009-04-30	94242
2	2009-07-31	100876
3	2009-10-31	99373
4	2010-01-31	113594

```
In [26]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56 entries, 0 to 55
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  56 non-null    datetime64[ns]
1   Quarterly Revenue     56 non-null    int64
dtypes: datetime64[ns](1), int64(1)
memory usage: 1.0 KB
```

```
In [27]: 1 # Setting Date as Index
2 df = df.set_index('Date')
3 df.head()
```

```
Out[27]:
```

	Quarterly Revenue
Date	
2009-01-31	108627
2009-04-30	94242
2009-07-31	100876
2009-10-31	99373
2010-01-31	113594


In [28]:

```
1 # Plotting the data
2 df.plot(figsize=(16,4),legend=True)
3 plt.title('Walmart Quarterly Revenue')
4 plt.show()
```



In [29]:

```
1 # Decomposition the data
2 from pylab import rcParams
3 rcParams['figure.figsize'] = 18, 8
4
5 decomposition = sm.tsa.seasonal_decompose(df, model = 'additive')
6 fig = decomposition.plot()
7 plt.show()
```




```
In [30]: 1 # Dividing the data into training and testing
2 # Plotting the data
3 import seaborn as sns
4 df['Date'] = df.index
5 train = df[df['Date'] < pd.to_datetime("2020-12", format='%Y-%m')]
6 train['train'] = train['Quarterly Revenue']
7 del train['Date']
8 del train['Quarterly Revenue']
9 test = df[df['Date'] >= pd.to_datetime("2020-12", format='%Y-%m')]
10 del test['Date']
11 test['test'] = test['Quarterly Revenue']
12 del test['Quarterly Revenue']
13 plt.plot(train, color = "black")
14 plt.plot(test, color = "red")
15 plt.title("Train/Test split for Quarterly Revenue")
16 plt.ylabel("Quarterly Revenue")
17 plt.xlabel('Date')
18 sns.set()
19 plt.show()
```

C:\Users\ravit\AppData\Local\Temp\ipykernel_14256\2575482981.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

train['train'] = train['Quarterly Revenue']
C:\Users\ravit\AppData\Local\Temp\ipykernel_14256\2575482981.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
test['test'] = test['Quarterly Revenue']



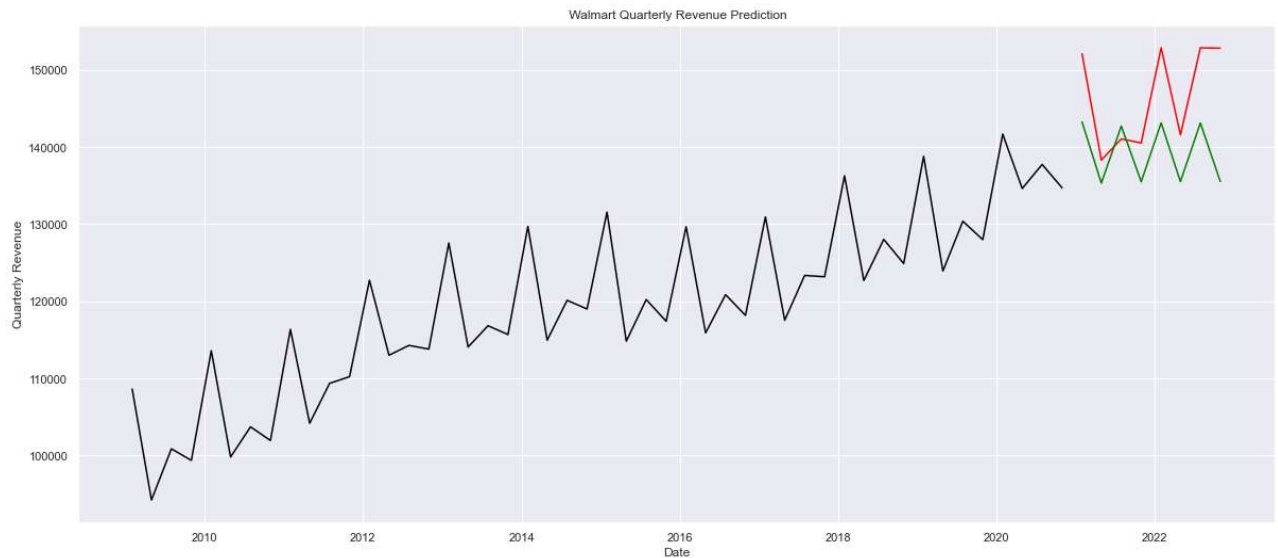
Arima Model

```
In [31]: 1 # Applying ARIMA Model
2 from pmdarima.arima import auto_arima
3 model = auto_arima(train, trace=True, error_action='ignore', suppress_warnings=True)
4 model.fit(train)
5 forecast = model.predict(n_periods=len(test))
6 forecast = pd.DataFrame(forecast, index = test.index, columns=['Prediction'])
```

```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=955.046, Time=0.20 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=1000.490, Time=0.01 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=987.989, Time=0.02 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=990.035, Time=0.04 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=998.653, Time=0.01 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.25 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=inf, Time=0.26 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.38 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=953.871, Time=0.27 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=952.720, Time=0.21 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=991.409, Time=0.07 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=952.424, Time=0.26 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=984.050, Time=0.08 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=952.498, Time=0.40 sec
ARIMA(1,1,5)(0,0,0)[0] intercept : AIC=950.471, Time=0.36 sec
ARIMA(0,1,5)(0,0,0)[0] intercept : AIC=980.089, Time=0.09 sec
ARIMA(2,1,5)(0,0,0)[0] intercept : AIC=951.493, Time=0.48 sec
ARIMA(1,1,5)(0,0,0)[0] : AIC=949.092, Time=0.30 sec
ARIMA(0,1,5)(0,0,0)[0] : AIC=978.866, Time=0.09 sec
ARIMA(1,1,4)(0,0,0)[0] : AIC=951.495, Time=0.26 sec
ARIMA(2,1,5)(0,0,0)[0] : AIC=953.488, Time=0.48 sec
ARIMA(0,1,4)(0,0,0)[0] : AIC=982.276, Time=0.08 sec
ARIMA(2,1,4)(0,0,0)[0] : AIC=951.561, Time=0.39 sec
```

```
Best model: ARIMA(1,1,5)(0,0,0)[0]
Total fit time: 5.039 seconds
```

```
In [32]: 1 # Plotting the prediction
2 plt.plot(train, color = "black")
3 plt.plot(test, color = "red")
4 plt.plot(forecast, color = "green")
5 plt.title(" Walmart Quarterly Revenue Prediction")
6 plt.ylabel("Quarterly Revenue")
7 plt.xlabel('Date')
8 sns.set()
9 plt.show()
```



```
In [33]: 1 from math import sqrt
2 from sklearn.metrics import mean_squared_error
3 rms = sqrt(mean_squared_error(test, forecast))
4 print("RMSE: ", rms)
```

```
RMSE: 8948.314010936627
```

SARIMA Model

```
In [34]: 1 df=pd.read_excel('Walmart Quarterly Revenue.xlsx')
2 df = df.set_index('Date')
```

```
In [35]: 1 # set the typical ranges for p, d, q
2 p = d = q = range(0, 2)
3
4 #take all possible combination for p, d and q
5 pdq = list(itertools.product(p, d, q))
6 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
7
8 print('Examples of parameter combinations for Seasonal ARIMA...')
9 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[1]))
10 print('SARIMAX: {} x {}'.format(pdq[1], seasonal_pdq[2]))
11 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[3]))
12 print('SARIMAX: {} x {}'.format(pdq[2], seasonal_pdq[4]))
```

Examples of parameter combinations for Seasonal ARIMA...

SARIMAX: (0, 0, 1) x (0, 0, 1, 12)

SARIMAX: (0, 0, 1) x (0, 1, 0, 12)

SARIMAX: (0, 1, 0) x (0, 1, 1, 12)

SARIMAX: (0, 1, 0) x (1, 0, 0, 12)

```
In [36]: 1 # Using Grid Search find the optimal set of parameters that yields the best performance
2 for param in pdq:
3     for param_seasonal in seasonal_pdq:
4         try:
5             mod = sm.tsa.statespace.SARIMAX(df, order = param, seasonal_order = param_seasonal, enforce_stationary = False, enforce_trend = False)
6             result = mod.fit()
7             print('SARIMA{}x{}12 - AIC:{}'.format(param, param_seasonal, result.aic))
8         except:
9             continue
```

SARIMA(0, 0, 0)x(0, 0, 0, 12)12 - AIC:1474.5164167195842

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

```
In [37]: 1 #Fitting the SARIMA model using above optimal combination of p, d, q (optimal means combination at which we got lowest AIC score)
2
3 model = sm.tsa.statespace.SARIMAX(df, order = (1, 1, 1),
4                                   seasonal_order = (1, 1, 0, 12)
5                                   )
6 result = model.fit()
7 print(result.summary().tables[1])
```

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)

R:\Anaconda\envs\general\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency Q-OCT will be used.

self._init_dates(dates, freq)


```
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1          -0.5736      0.586     -0.978     0.328     -1.723     0.576
ma.L1           0.4940      0.632      0.781     0.435     -0.745     1.733
ar.S.L12        0.0019      0.164      0.011     0.991     -0.320     0.324
sigma2         6.155e+06    2.96e-08    2.08e+14    0.000    6.15e+06    6.15e+06
=====
```

```
In [38]: 1 prediction = result.get_prediction(start = pd.to_datetime('2021-04-30'), dynamic = False)
2 prediction_ci = prediction.conf_int()
3 prediction_ci
```

Out[38]:

	lower Quarterly Revenue	upper Quarterly Revenue
2021-04-30	133299.243433	143023.993306
2021-07-31	138972.339021	148697.088894
2021-10-31	133165.550097	142890.299970
2022-01-31	149300.397194	159025.147067
2022-04-30	133392.899885	143117.649758
2022-07-31	142751.969848	152476.719722
2022-10-31	145428.832133	155153.582006


```
In [39]: 1 #Visualize the forecasting
2 ax = df['2009:'].plot(label = 'observed')
3 prediction.predicted_mean.plot(ax = ax, label = 'One-step ahead Forecast', alpha = 0.7, figsize = (14, 7))
4 ax.fill_between(prediction_ci.index, prediction_ci.iloc[:, 0], prediction_ci.iloc[:, 1], color = 'k', alpha = 0.2)
5 ax.set_xlabel("year")
6 ax.set_ylabel('Quarterly Revenue')
7 plt.legend()
8 plt.show()
```



```
In [40]: 1 # Evaluation metrics are Squared Mean Error(SME) and Root Mean Squared Error(RMSE)
2
3 from sklearn.metrics import mean_squared_error
4
5 y_hat = prediction.predicted_mean
6 y_truth = df['2021-04-30:']
7 mse = mean_squared_error(y_truth,y_hat)
8 rmse = np.sqrt(mse)
9
10 print('The Mean Squared Error of our forecasts is', mse)
11 print('The Root Mean Squared Error of our forecasts is', rmse)
```

The Mean Squared Error of our forecasts is 8648353.268182157
The Root Mean Squared Error of our forecasts is 2940.808267837629

```
In [41]: 1 # forecasting for out of sample data
2 pred_uc = result.get_forecast(steps = 10)
3 pred_ci = pred_uc.conf_int()
4
5 ax = df.plot(label = 'observed', figsize = (14, 7))
6 pred_uc.predicted_mean.plot(ax = ax, label = 'forecast')
7 ax.fill_between(pred_ci.index, pred_ci.iloc[:, 0], pred_ci.iloc[:, 1], color = 'k', alpha = 0.25)
8 ax.set_xlabel('Date')
9 ax.set_ylabel('Quarterly Revenue')
10
11 plt.legend()
12 plt.show()
13
```



DNN MODEL

```
In [42]: 1 def convert2matrix(data_arr, look_back):
2         X, Y = [], []
3         for i in range(len(data_arr)-look_back):
4             d=i+look_back
5             X.append(data_arr[i:d,0])
6             Y.append(data_arr[d,0])
7         return np.array(X).astype('int'), np.array(Y).astype('int')
```

```
In [43]: 1 df=pd.read_excel('Walmart Quarterly Revenue.xlsx')
2
3 df = df.set_index('Date')
4
```

```
In [44]: 1 df1 = df
2 #Split data set into testing dataset and train dataset
3 train_size = 49
4 train, test =df1.values[0:train_size,:],df1.values[train_size:len(df1.values),:]
5 # setup look_back window
6 look_back = 4
7 #convert dataset into right shape in order to input into the DNN
8 trainX, trainY = convert2matrix(train, look_back)
9 testX, testY = convert2matrix(test, look_back)
```

```
In [45]: 1 from keras.models import Sequential
2 from keras.layers import Dense
3 def model_dnn(look_back):
4     model=Sequential()
5     model.add(Dense(units=32, input_dim=look_back, activation='relu'))
6     model.add(Dense(8, activation='relu'))
7     model.add(Dense(1))
8     model.compile(loss='mean_squared_error', optimizer='adam',metrics = ['mse', 'mae'])
9     return model
```

```
In [46]: 1 model=model_dnn(look_back)
2 history=model.fit(trainX,trainY, epochs=500, batch_size=4, verbose=1, validation_data=(testX,testY),shuffle=False)
```

```
Epoch 1/500
12/12 [=====] - 1s 16ms/step - loss: 17379080192.0000 - mse: 17379080192.0000 - mae: 131621.1875 - val_loss: 21999052800.0000 - val_mse: 21999052800.0000 - val_mae: 148240.0156
Epoch 2/500
12/12 [=====] - 0s 3ms/step - loss: 12832726016.0000 - mse: 12832726016.0000 - mae: 113052.1875 - val_loss: 15932937216.0000 - val_mse: 15932937216.0000 - val_mae: 126136.8750
Epoch 3/500
12/12 [=====] - 0s 3ms/step - loss: 9045407744.0000 - mse: 9045407744.0000 - mae: 94868.5781 - val_loss: 10838700032.0000 - val_mse: 10838700032.0000 - val_mae: 104005.6172
Epoch 4/500
12/12 [=====] - 0s 3ms/step - loss: 5943562752.0000 - mse: 5943562752.0000 - mae: 76815.9219 - val_loss: 6807878144.0000 - val_mse: 6807878144.0000 - val_mae: 82382.6172
Epoch 5/500
12/12 [=====] - 0s 3ms/step - loss: 3581512448.0000 - mse: 3581512448.0000 - mae: 59482.8984 - val_loss: 3852875776.0000 - val_mse: 3852875776.0000 - val_mae: 61911.6406
Epoch 6/500
12/12 [=====] - 0s 3ms/step - loss: 1903954176.0000 - mse: 1903954176.0000 - mae: 43107.1484 - val_loss: 1883590656.0000 - val_mse: 1883590656.0000 - val_mae: 43177.4336
Epoch 7/500
12/12 [=====] - 0s 3ms/step - loss: 964573264.0000 - mse: 964573264.0000 - mae: 28565.4600 - val_loss: 964573264.0000 - val_mse: 964573264.0000 - val_mae: 28565.4600
```


```
In [47]: 1 def model_loss(history):
2     plt.figure(figsize=(8,4))
3     plt.plot(history.history['loss'], label='Train Loss')
4     plt.plot(history.history['val_loss'], label='Test Loss')
5     plt.title('model loss')
6     plt.ylabel('loss')
7     plt.xlabel('epochs')
8     plt.legend(loc='upper right')
9     plt.show();
```

```
In [48]: 1 train_score = model.evaluate(trainX, trainY, verbose=0)
2 print('Train Root Mean Squared Error(RMSE): %.2f; Train Mean Absolute Error(MAE) : %.2f '
3 % (np.sqrt(train_score[1]), train_score[2]))
4 test_score = model.evaluate(testX, testY, verbose=0)
5 print(train_score)
6 print(test_score)
7 print('Test Root Mean Squared Error(RMSE): %.2f; Test Mean Absolute Error(MAE) : %.2f '
8 % (np.sqrt(test_score[1]), test_score[2]))
```

```
Train Root Mean Squared Error(RMSE): 2998.15; Train Mean Absolute Error(MAE) : 2389.67
[8988885.0, 8988885.0, 2389.6708984375]
[56950688.0, 56950688.0, 6428.11962890625]
Test Root Mean Squared Error(RMSE): 7546.57; Test Mean Absolute Error(MAE) : 6428.12
```

```
In [49]: 1 def prediction_plot(testY, test_predict):
2         len_prediction=[x for x in range(len(testY))]
3         plt.figure(figsize=(8,4))
4         plt.plot(len_prediction, testY[:8], marker='.', label="actual")
5         plt.plot(len_prediction, test_predict[:8], 'r', label="prediction")
6         plt.tight_layout()
7         sns.despine(top=True)
8         plt.subplots_adjust(left=0.07)
9         plt.ylabel('Ads Daily Spend', size=15)
10        plt.xlabel('Time step', size=15)
11        plt.legend(fontsize=15)
12        plt.show();
```

```
In [50]: 1 test_predict = model.predict(testX)
2         prediction_plot(testY, test_predict)
```



GRU and BiLSTM Models

```
In [51]: 1 df=pd.read_excel('Walmart Quarterly Revenue.xlsx')
2         df.head()
```

```
Out[51]:
```

	Date	Quarterly Revenue
0	2009-01-31	108627
1	2009-04-30	94242
2	2009-07-31	100876
3	2009-10-31	99373
4	2010-01-31	113594

```
In [52]: 1 df = df.set_index('Date')
2         df.head()
```

```
Out[52]:
```

	Date	Quarterly Revenue
	2009-01-31	108627
	2009-04-30	94242
	2009-07-31	100876
	2009-10-31	99373
	2010-01-31	113594

```
In [53]: 1 # Split train data and test data
2         train_size = int(len(df)*0.8)
3
4         train_data = df.iloc[:train_size]
5         test_data = df.iloc[train_size:]
```

```
In [54]: 1 scaler = MinMaxScaler().fit(train_data)
2         train_scaled = scaler.transform(train_data)
3         test_scaled = scaler.transform(test_data)
```



```
In [55]: 1 # Create input dataset
2 def create_dataset (X, look_back = 1):
3     Xs, ys = [], []
4
5     for i in range(len(X)-look_back):
6         v = X[i:i+look_back]
7         Xs.append(v)
8         ys.append(X[i+look_back])
9
10    return np.array(Xs), np.array(ys)
11 LOOK_BACK = 4
12 X_train, y_train = create_dataset(train_scaled,LOOK_BACK)
13 X_test, y_test = create_dataset(test_scaled,LOOK_BACK)
14 # Print data shape
15 print('X_train.shape: ', X_train.shape)
16 print('y_train.shape: ', y_train.shape)
17 print('X_test.shape: ', X_test.shape)
18 print('y_test.shape: ', y_test.shape)
```

```
X_train.shape: (40, 4, 1)
y_train.shape: (40, 1)
X_test.shape: (8, 4, 1)
y_test.shape: (8, 1)
```

```
In [56]: 1 # Create BiLSTM model
2 def create_bilstm(units):
3     model = Sequential()
4     # Input Layer
5     model.add(Bidirectional(
6         LSTM(units = units, return_sequences=True),
7         input_shape=(X_train.shape[1], X_train.shape[2])))
8     # Hidden Layer
9     model.add(Bidirectional(LSTM(units = units)))
10    model.add(Dense(1))
11    #Compile model
12    model.compile(optimizer='adam',loss='mse')
13    return model
14 model_bilstm = create_bilstm(64)
15 # Create GRU model
16 def create_gru(units):
17     model = Sequential()
18     # Input Layer
19     model.add(GRU (units = units, return_sequences = True,
20         input_shape = [X_train.shape[1], X_train.shape[2]]))
21     model.add(Dropout(0.2))
22     # Hidden Layer
23     model.add(GRU(units = units))
24     model.add(Dropout(0.2))
25     model.add(Dense(units = 1))
26     #Compile model
27     model.compile(optimizer='adam',loss='mse')
28     return model
29 model_gru = create_gru(64)
```

In [57]:

```
1 def fit_model(model):
2     early_stop = keras.callbacks.EarlyStopping(monitor = 'val_loss',
3                                                 patience = 10)
4     history = model.fit(X_train, y_train, epochs = 100,
5                         validation_split = 0.2,
6                         batch_size = 16, shuffle = False,
7                         callbacks = [early_stop])
8     return history
9 history_gru = fit_model(model_gru)
```



```
Epoch 1/100
2/2 [=====] - 3s 604ms/step - loss: 0.2240 - val_loss: 0.3596
Epoch 2/100
2/2 [=====] - 0s 28ms/step - loss: 0.1307 - val_loss: 0.2115
Epoch 3/100
2/2 [=====] - 0s 27ms/step - loss: 0.0726 - val_loss: 0.1028
Epoch 4/100
2/2 [=====] - 0s 27ms/step - loss: 0.0423 - val_loss: 0.0369
Epoch 5/100
2/2 [=====] - 0s 26ms/step - loss: 0.0261 - val_loss: 0.0171
Epoch 6/100
2/2 [=====] - 0s 26ms/step - loss: 0.0322 - val_loss: 0.0242
Epoch 7/100
2/2 [=====] - 0s 30ms/step - loss: 0.0418 - val_loss: 0.0271
Epoch 8/100
2/2 [=====] - 0s 25ms/step - loss: 0.0453 - val_loss: 0.0213
Epoch 9/100
2/2 [=====] - 0s 24ms/step - loss: 0.0284 - val_loss: 0.0175
Epoch 10/100
2/2 [=====] - 0s 26ms/step - loss: 0.0182 - val_loss: 0.0204
Epoch 11/100
2/2 [=====] - 0s 29ms/step - loss: 0.0209 - val_loss: 0.0283
Epoch 12/100
2/2 [=====] - 0s 26ms/step - loss: 0.0186 - val_loss: 0.0372
Epoch 13/100
2/2 [=====] - 0s 27ms/step - loss: 0.0196 - val_loss: 0.0432
Epoch 14/100
2/2 [=====] - 0s 26ms/step - loss: 0.0233 - val_loss: 0.0445
Epoch 15/100
2/2 [=====] - 0s 28ms/step - loss: 0.0242 - val_loss: 0.0411
```

In [58]:

```
1 history_bilstm = fit_model(model_bilstm)
```

```
Epoch 1/100
2/2 [=====] - 6s 1s/step - loss: 0.2862 - val_loss: 0.5206
Epoch 2/100
2/2 [=====] - 0s 29ms/step - loss: 0.2158 - val_loss: 0.3906
Epoch 3/100
2/2 [=====] - 0s 29ms/step - loss: 0.1541 - val_loss: 0.2712
Epoch 4/100
2/2 [=====] - 0s 34ms/step - loss: 0.0992 - val_loss: 0.1632
Epoch 5/100
2/2 [=====] - 0s 27ms/step - loss: 0.0531 - val_loss: 0.0751
Epoch 6/100
2/2 [=====] - 0s 30ms/step - loss: 0.0227 - val_loss: 0.0236
Epoch 7/100
2/2 [=====] - 0s 28ms/step - loss: 0.0170 - val_loss: 0.0206
Epoch 8/100
2/2 [=====] - 0s 29ms/step - loss: 0.0329 - val_loss: 0.0349
Epoch 9/100
2/2 [=====] - 0s 28ms/step - loss: 0.0420 - val_loss: 0.0304
Epoch 10/100
2/2 [=====] - 0s 31ms/step - loss: 0.0341 - val_loss: 0.0190
Epoch 11/100
2/2 [=====] - 0s 30ms/step - loss: 0.0227 - val_loss: 0.0173
Epoch 12/100
2/2 [=====] - 0s 32ms/step - loss: 0.0169 - val_loss: 0.0260
Epoch 13/100
2/2 [=====] - 0s 30ms/step - loss: 0.0170 - val_loss: 0.0380
Epoch 14/100
2/2 [=====] - 0s 35ms/step - loss: 0.0195 - val_loss: 0.0468
Epoch 15/100
2/2 [=====] - 0s 30ms/step - loss: 0.0214 - val_loss: 0.0497
Epoch 16/100
2/2 [=====] - 0s 32ms/step - loss: 0.0216 - val_loss: 0.0468
Epoch 17/100
2/2 [=====] - 0s 33ms/step - loss: 0.0201 - val_loss: 0.0399
Epoch 18/100
2/2 [=====] - 0s 31ms/step - loss: 0.0179 - val_loss: 0.0317
Epoch 19/100
2/2 [=====] - 0s 30ms/step - loss: 0.0161 - val_loss: 0.0245
Epoch 20/100
2/2 [=====] - 0s 30ms/step - loss: 0.0154 - val_loss: 0.0198
Epoch 21/100
2/2 [=====] - 0s 29ms/step - loss: 0.0158 - val_loss: 0.0175
```

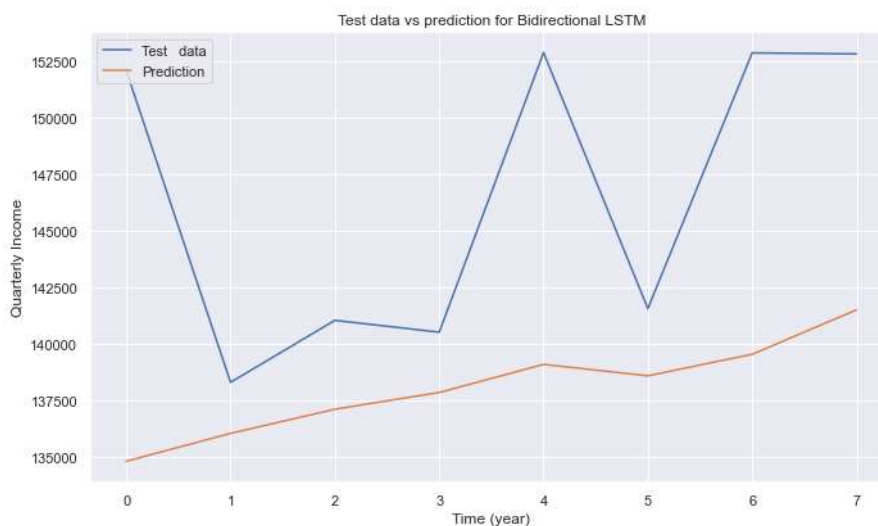
```
In [59]: 1 def plot_loss(history, model_name):
2         plt.figure(figsize = (10, 6))
3         plt.plot(history.history['loss'])
4         plt.plot(history.history['val_loss'])
5         plt.title('Model Train vs Validation Loss for ' + model_name)
6         plt.ylabel('Loss')
7         plt.xlabel('epoch')
8         plt.legend(['Train loss', 'Validation loss'], loc='upper right')
9
10        plot_loss(history_gru, 'GRU')
11        plot_loss(history_bilstm, 'Bidirectional LSTM')
```



```

In [60]: 1 # Make prediction
2 def prediction(model):
3     prediction = model.predict(X_test)
4     prediction = scaler.inverse_transform(prediction)
5     return prediction
6 prediction_gru = prediction(model_gru)
7 prediction_bilstm = prediction(model_bilstm)
8 # Plot test data vs prediction
9 def plot_future(prediction, model_name, y_test):
10    plt.figure(figsize=(10, 6))
11    range_future = len(prediction)
12    plt.plot(np.arange(range_future), np.array(scaler.inverse_transform(y_test)),
13             label='Test data')
14    plt.plot(np.arange(range_future),
15             np.array(prediction), label='Prediction')
16    plt.title('Test data vs prediction for ' + model_name)
17    plt.legend(loc='upper left')
18    plt.xlabel('Time (year)')
19    plt.ylabel('Quarterly Income')
20
21 plot_future(prediction_gru, 'GRU', y_test)
22 plot_future(prediction_bilstm, 'Bidirectional LSTM', y_test)

```



```

In [61]: 1 def evaluate_prediction(predictions, actual, model_name):
2     errors = predictions - actual
3     mse = np.square(errors).mean()
4     rmse = np.sqrt(mse)
5     mae = np.abs(errors).mean()
6     print(model_name + ':')
7     print('Mean Absolute Error: {:.4f}'.format(mae))
8     print('Root Mean Square Error: {:.4f}'.format(rmse))
9     print('')
10 evaluate_prediction(prediction_gru, scaler.inverse_transform(y_test), 'GRU')
11 evaluate_prediction(prediction_bilstm, scaler.inverse_transform(y_test), 'Bidirectional LSTM')

```

GRU:
Mean Absolute Error: 15263.1260
Root Mean Square Error: 16381.6322

Bidirectional LSTM:
Mean Absolute Error: 8430.6992
Root Mean Square Error: 10176.7920

```
In [62]: 1 def evaluate_prediction(predictions, actual, model_name):
2         errors = predictions - actual
3         mse = np.square(errors).mean()
4         rmse = np.sqrt(mse)
5         mae = np.abs(errors).mean()
6         print(model_name + ':')
7         print('Mean Absolute Error: {:.4f}'.format(mae))
8         print('Root Mean Square Error: {:.4f}'.format(rmse))
9         print('')
10        evaluate_prediction(scaler.transform(prediction_gru), y_test, 'GRU')
11        evaluate_prediction(scaler.transform(prediction_bilstm), y_test, 'Bidirectional LSTM')
```

GRU:

Mean Absolute Error: 0.3426

Root Mean Square Error: 0.3677

Bidirectional LSTM:

Mean Absolute Error: 0.1892

Root Mean Square Error: 0.2284

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

R:\Anaconda\envs\general\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names
warnings.warn(

In []:

1